Lecture 0: Introduction to Applied Research in Health Data Science

CSCI6410/4148 & EPAH6410

Finlay Maguire (finlay.maguire@dal.ca)
TA: Ehsan Baratnezhad (ethan.b@dal.ca)

Why am I teaching this course?

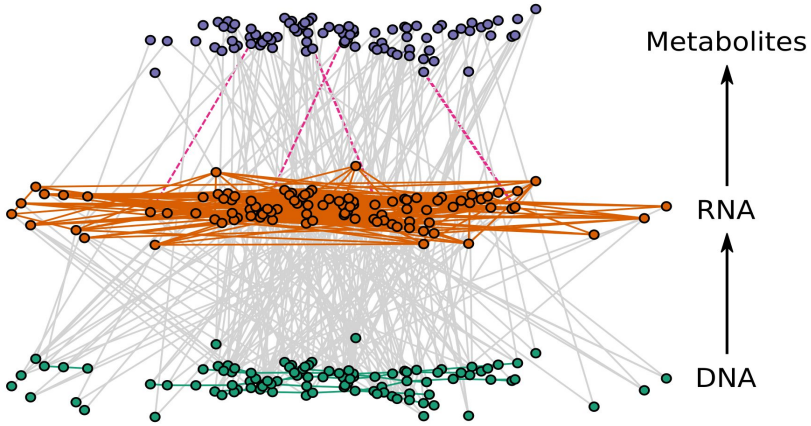

- **PhD (Bioinformatics)**: using large noisy datasets to understand how microbial systems and mechanisms evolve.

Why am I teaching this course?

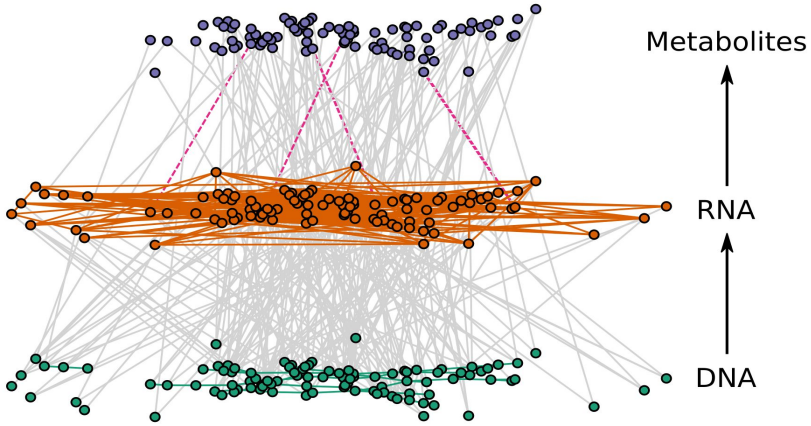

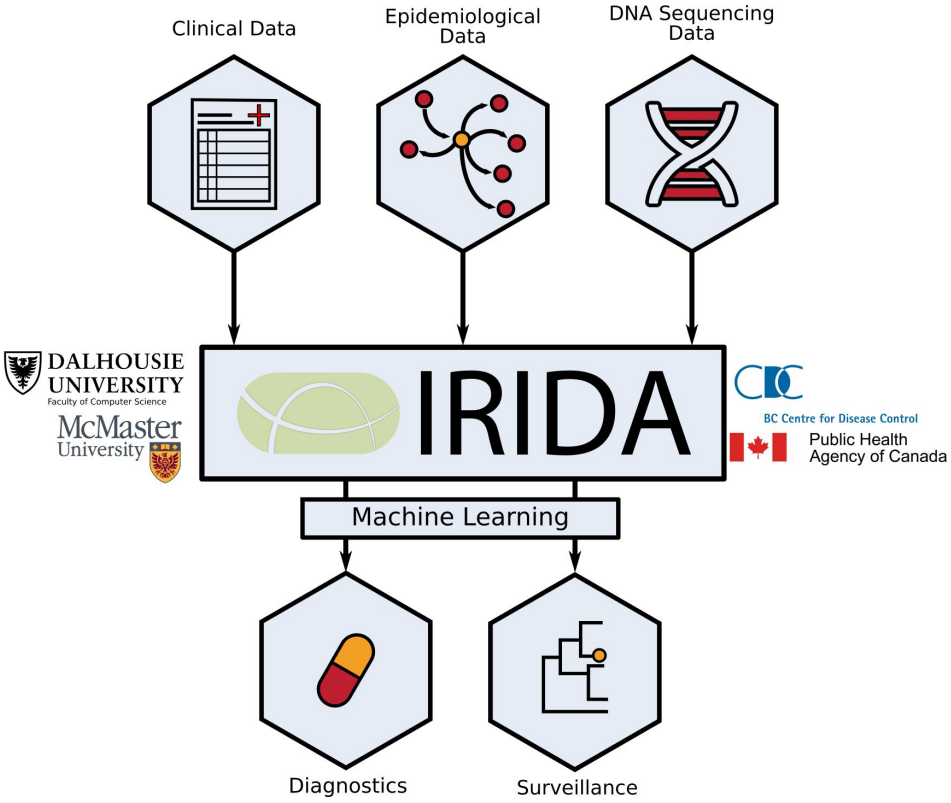

- **PhD (Bioinformatics)**: using large noisy datasets to understand how microbial systems and mechanisms evolve.
- **Postdoc (Genomic Epidemiology)**: using large noisy datasets to better diagnose, track and predict infectious diseases.

Why am I teaching this course?

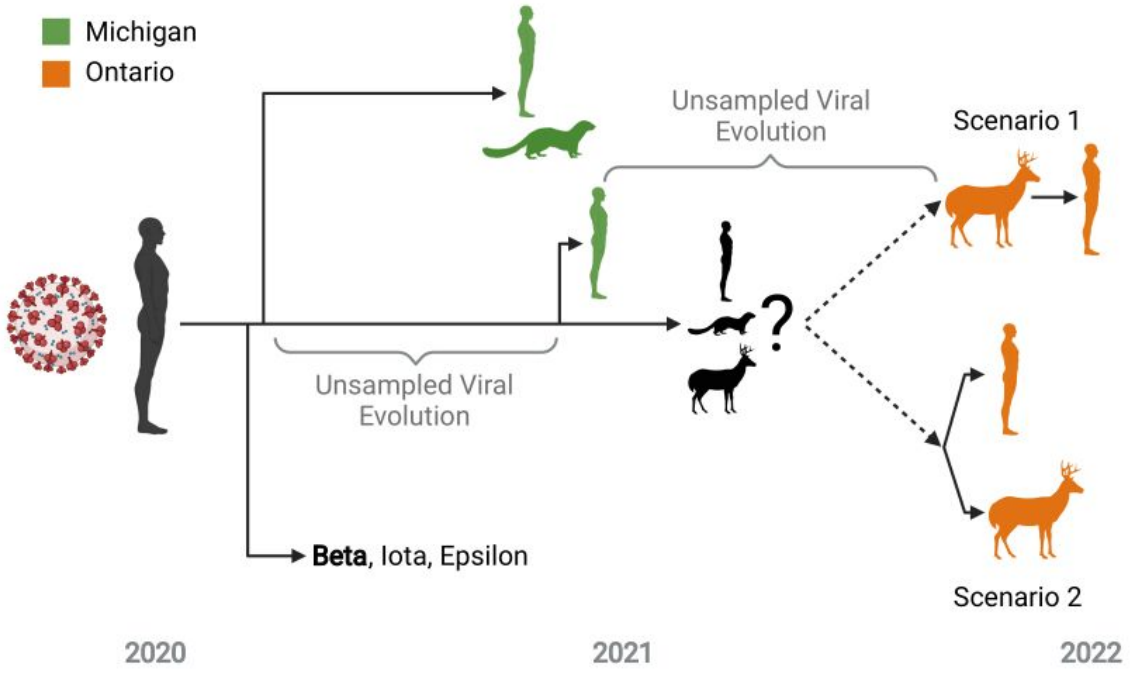

- **Research group**: using large noisy datasets:
  - Genomic epidemiology of infectious disease: **SARS-CoV-2**, **AMR**

Why am I teaching this course?

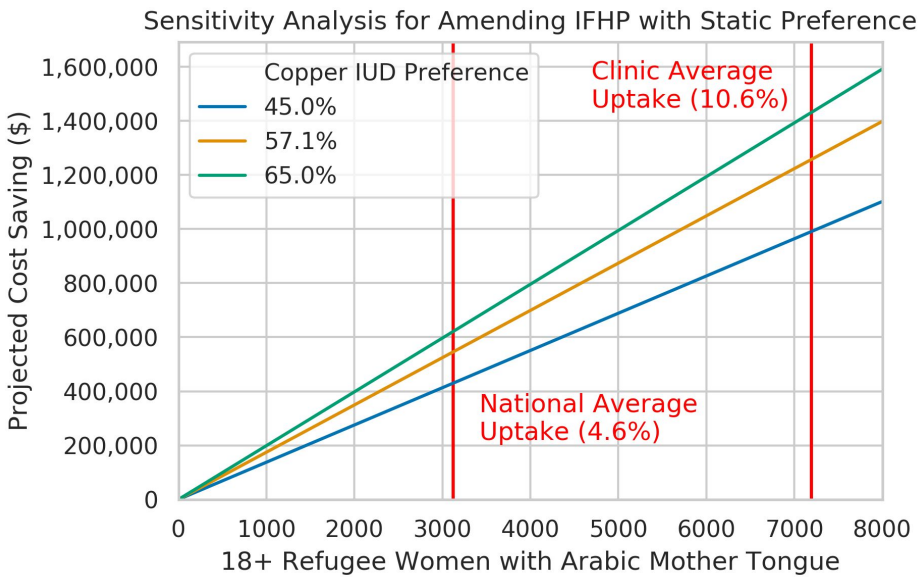

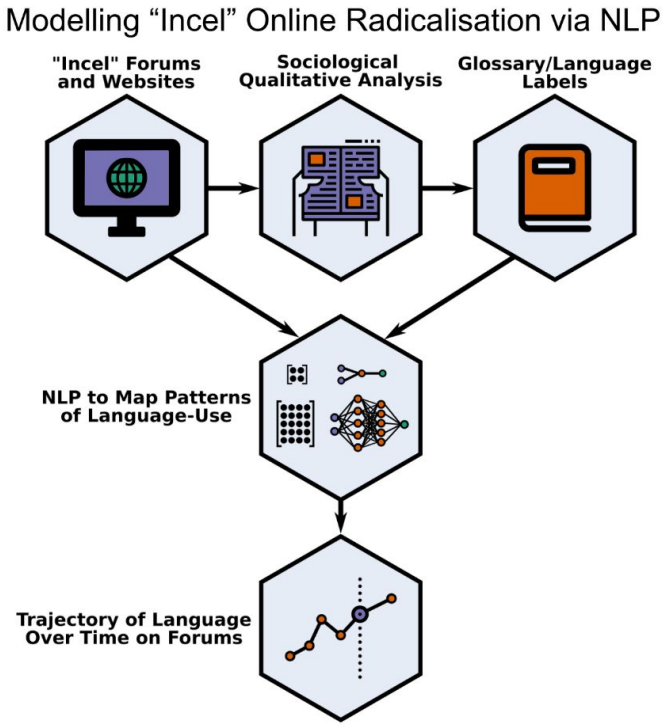

- **Research group**: using large noisy datasets:
  - Genomic epidemiology of infectious disease: **SARS-CoV-2**, **AMR**
  - Collaborations on socially/health focused problems: **refugee health**, **incel radicalisation, health inequality**

Overview of course

Applied Research in Health Data Science

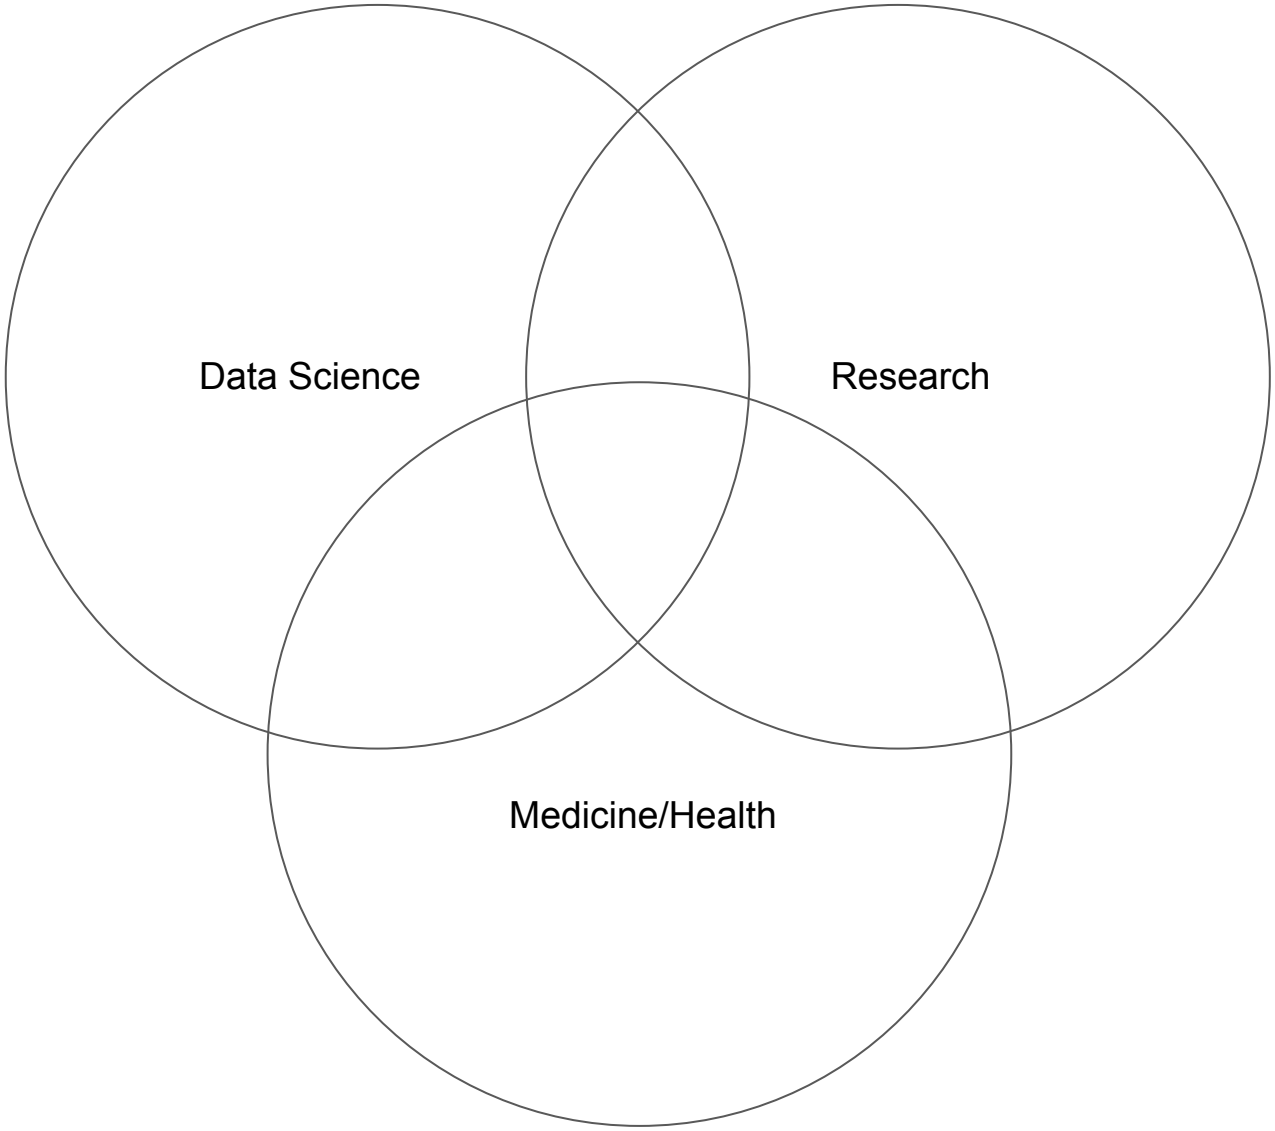

Applied Research in Health Data Science

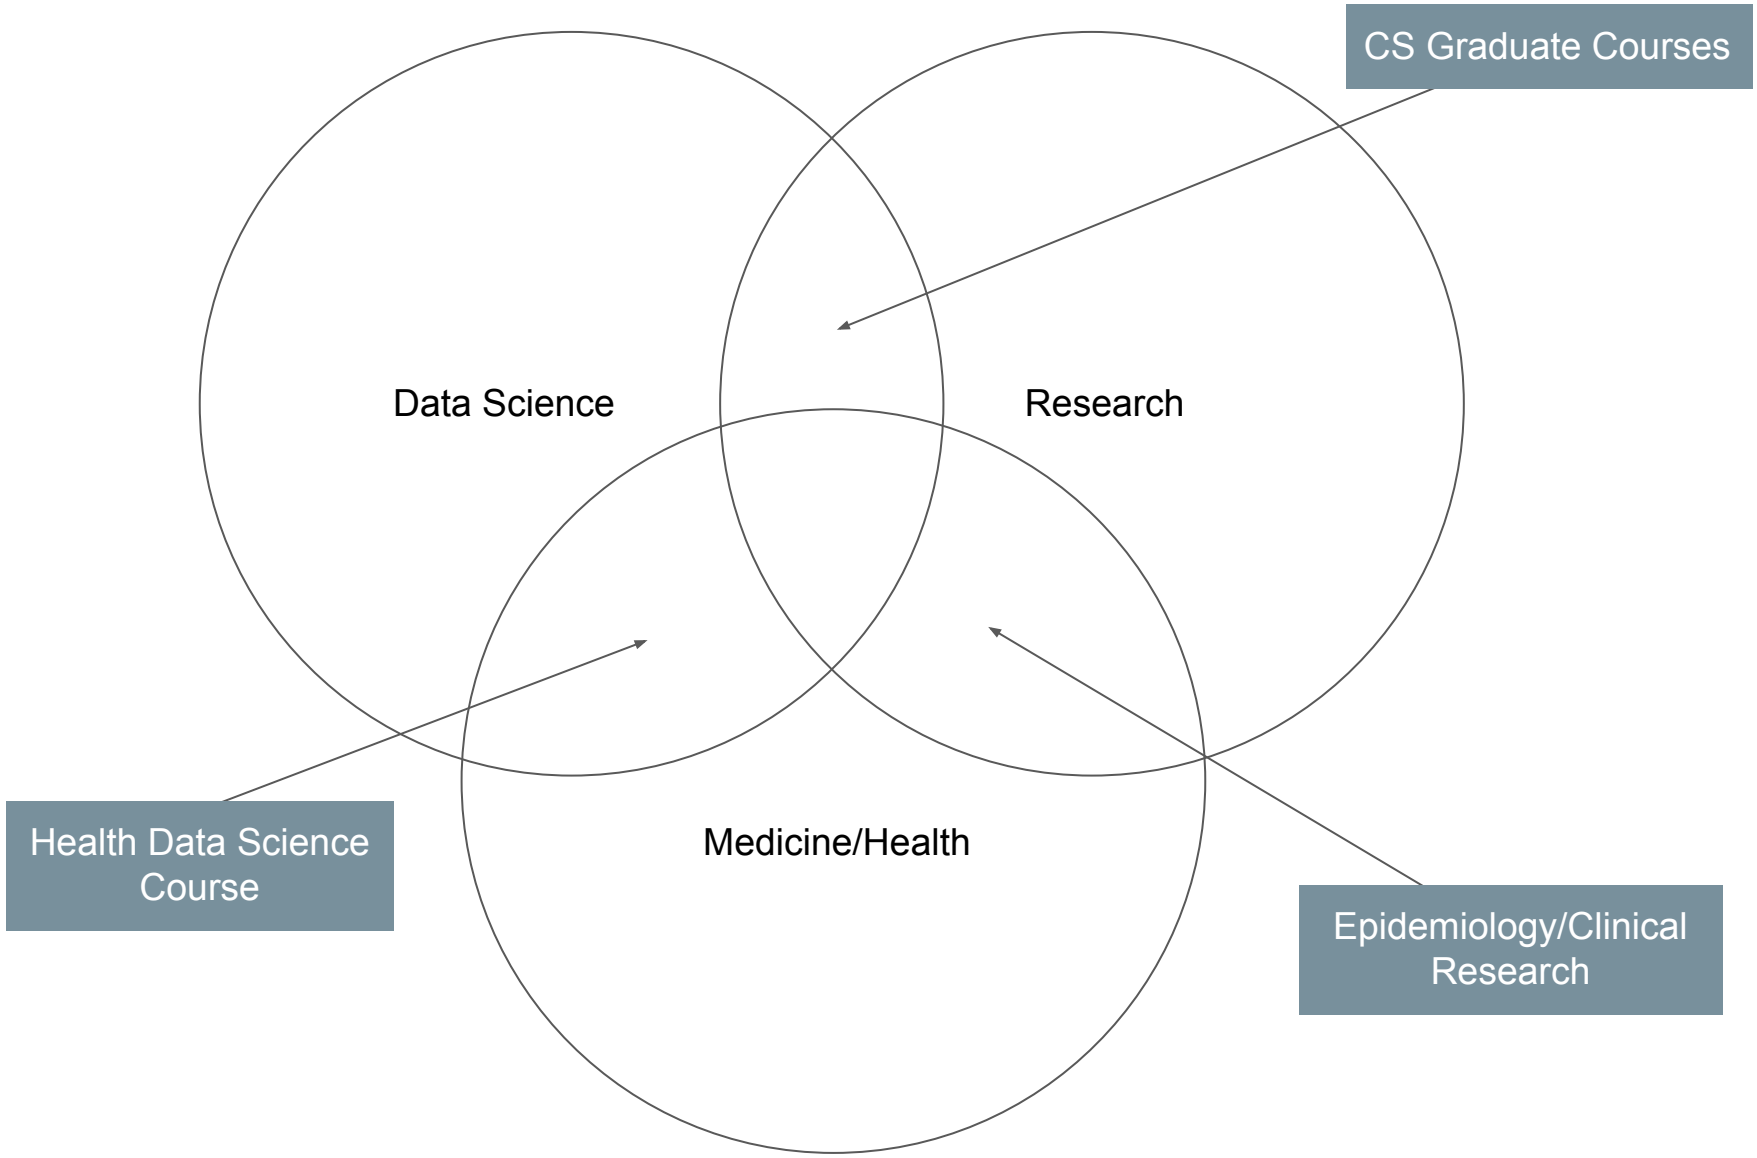

Applied Research in Health Data Science

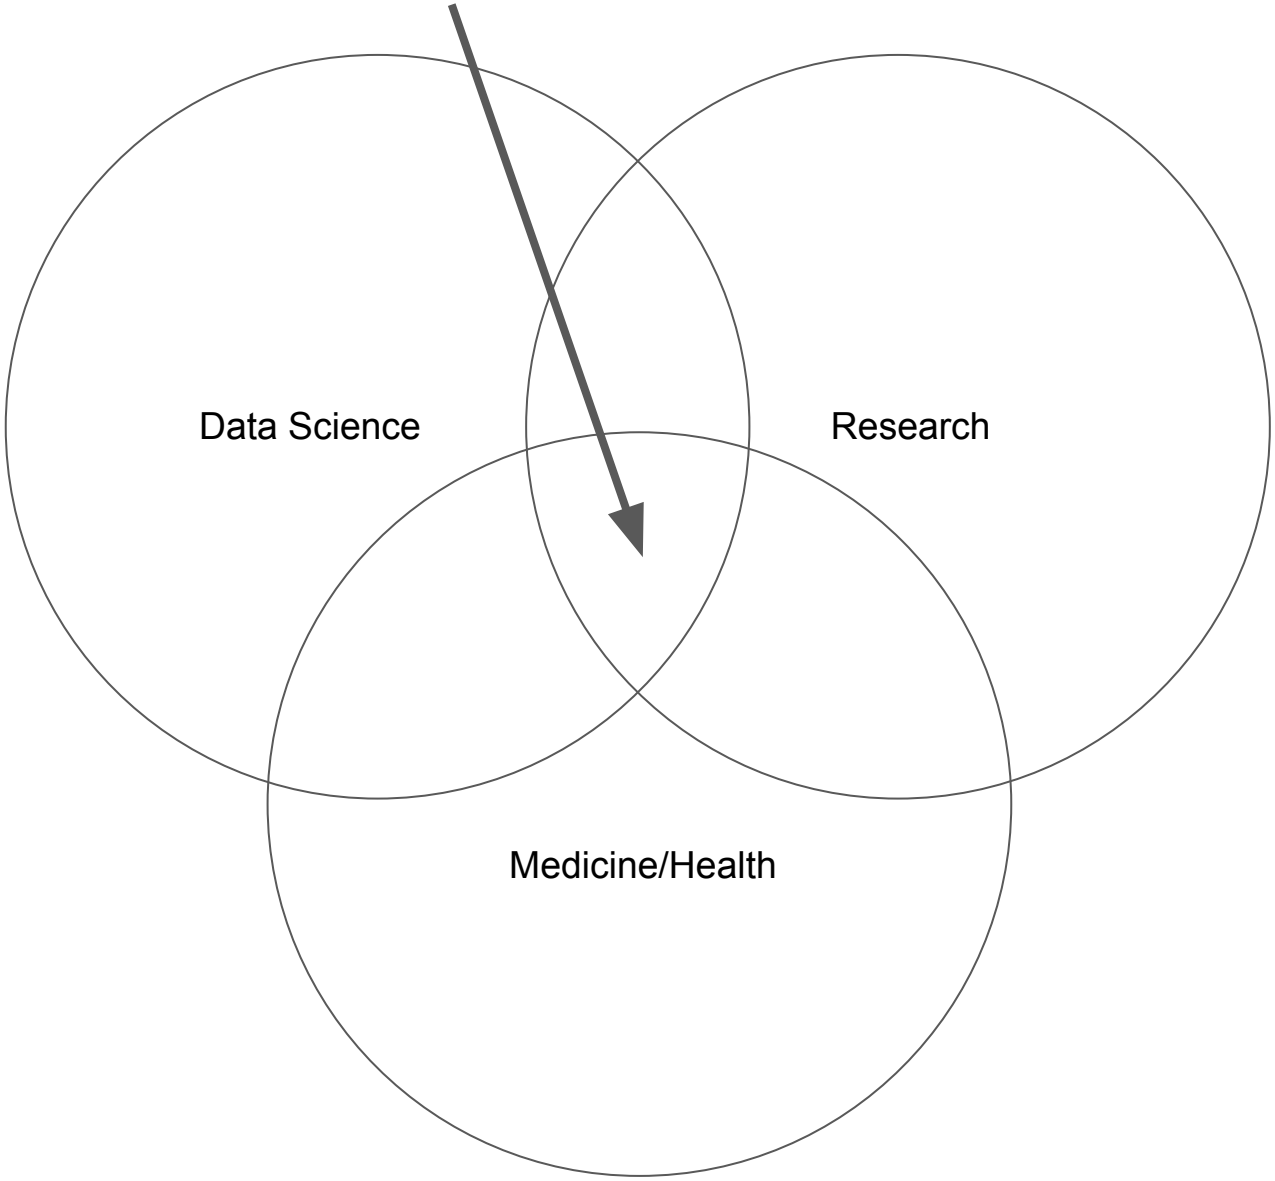

Learning Outcomes

1. Understand the **4 principal sources and data types** of medical data:
    a. longitudinal databases (tabular)
    b. electronic medical records (structured, semi-structured, and unstructured text)
    c. radiological imaging (image)
    d. physiological (signal and time-series).

Learning Outcomes

1. Understand the **4 principal sources and data types** of medical data:
   a. longitudinal databases (tabular)
   b. electronic medical records (structured, semi-structured, and unstructured text)
   c. radiological imaging (image)
   d. physiological (signal and time-series).
2. Identify and apply **appropriate type of method** to the analysis of each data type

Learning Outcomes

1. Understand the **4 principal sources and data types** of medical data:
   a. longitudinal databases (tabular)
   b. electronic medical records (structured, semi-structured, and unstructured text)
   c. radiological imaging (image)
   d. physiological (signal and time-series).
2. Identify and apply **appropriate type of method** to the analysis of each data type
3. Gain the technical skills necessary for effective health data science research including **data management, reproducibility**, and version control.

Learning Outcomes

1. Understand the **4 principal sources and data types** of medical data:
    a. longitudinal databases (tabular)
    b. electronic medical records (structured, semi-structured, and unstructured text)
    c. radiological imaging (image)
    d. physiological (signal and time-series).
2. Identify and apply **appropriate type of method** to the analysis of each data type
3. Gain the technical skills necessary for effective health data science research including **data management, reproducibility**, and version control.
4. Understand the key **collaborative, legal, ethical, and knowledge translation** concepts required in interdisciplinary health data science research.

Learning Outcomes

1. Understand the **4 principal sources and data types** of medical data:
   a. longitudinal databases (tabular)
   b. electronic medical records (structured, semi-structured, and unstructured text)
   c. radiological imaging (image)
   d. physiological (signal and time-series).
2. Identify and apply **appropriate type of method** to the analysis of each data type
3. Gain the technical skills necessary for effective health data science research including **data management, reproducibility**, and version control.
4. Understand the key **collaborative, legal, ethical, and knowledge translation** concepts required in interdisciplinary health data science research.
5. Critically **appraise research literature** in health data science.

Learning Outcomes

1. Understand the **4 principal sources and data types** of medical data:
    a. longitudinal databases (tabular)
    b. electronic medical records (structured, semi-structured, and unstructured text)
    c. radiological imaging (image)
    d. physiological (signal and time-series).
2. Identify and apply **appropriate type of method** to the analysis of each data type
3. Gain the technical skills necessary for effective health data science research including **data management, reproducibility**, and version control.
4. Understand the key **collaborative, legal, ethical, and knowledge translation** concepts required in interdisciplinary health data science research.
5. Critically **appraise research literature** in health data science.
6. Combine these skills to develop high-quality collaborative health data science **research proposals**

What is not covered in this course

- **Breadth/depth** of each data science method: *each could be multiple graduate CS courses*

What is not covered in this course

- **Breadth/depth** of each data science method: *each could be multiple graduate CS courses*
- **Breadth/depth** of medical research: *again could be a whole PhD program*

What is not covered in this course

- **Breadth/depth** of each data science method: *each could be multiple graduate CS courses*
- **Breadth/depth** of medical research: *again could be a whole PhD program*
- True **messiness** of real data: *provide tools but experience is invaluable*

What is not covered in this course

- **Breadth/depth** of each data science method: *each could be multiple graduate CS courses*
- **Breadth/depth** of medical research: *again could be a whole PhD program*
- True **messiness** of real data: *provide tools but experience is invaluable*
- Some important forms of medical data (e.g., genomics): *see CSCI4181/6810, EPAH6052 (partially), come speak to me if interested in this specifically.*

Course Structure

Overview of data types & analysis methods:

- **Lectures** (Monday/Wednesday)

Longitudinal Participation: (10%)

Course Structure

Overview of data types & analysis methods:

- **Lectures** (Monday/Wednesday)
- **Practical Exercises** (Friday/Monday)

Assessment: Submission of Practical Exercise Due the day before **following practical** (10% x 4)

(CSCI4148: drop lowest scoring assignment)

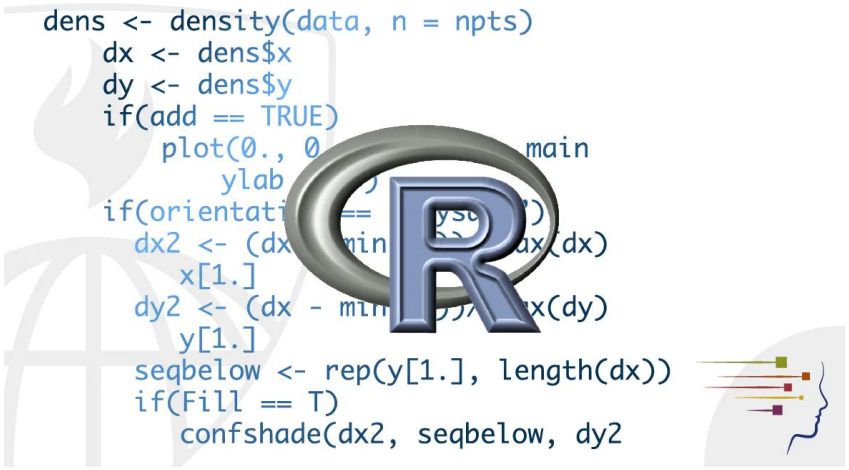

https://www.coursera.org/learn/r-programming

Longitudinal Participation: (10%)

Course Structure

Overview of data types & analysis methods:

- **Lectures** (Monday/Wednesday)
- **Practical Exercises** (Friday/Monday)

Assessment: Submission of Practical Exercise Due the day before **following practical** (10% x 4)

(CSCI4148: drop lowest scoring assignment)

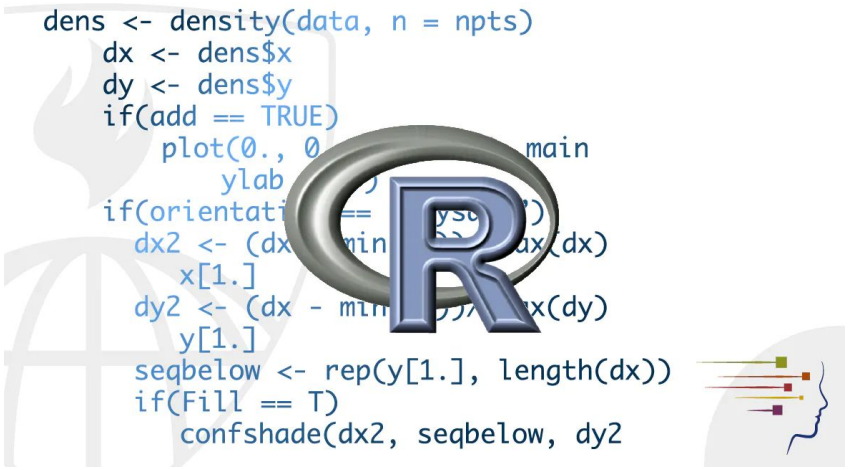

https://www.coursera.org/learn/r-programming

Research in health data science:

- **Journal Club** (Wednesday/Friday)

2 papers per week, randomly assigned rota for leading discussion of paper with rest of class.

Assessment:

Paper presentation (15%)

Participation in discussion (10%)

Longitudinal Participation: (10%)

Course Structure

Overview of data types & analysis methods:

- **Lectures** (Monday/Wednesday)
- **Practical Exercises** (Friday/Monday)

Assessment: Submission of Practical Exercise Due the day before **following practical** (10% x 4)

(CSCI4148: drop lowest scoring assignment)

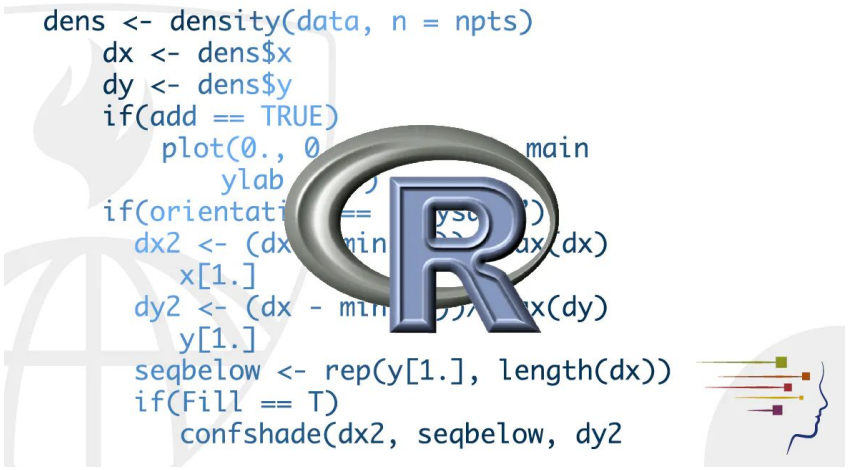

https://www.coursera.org/learn/r-programming

Research in health data science:

- **Journal Club** (Wednesday/Friday)

2 papers per week, randomly assigned rota for leading discussion of paper with rest of class.

Assessment:

Paper presentation (15%)

Participation in discussion (10%)

Development of a research proposal:

- **Class** (Wednesday/Friday)

Assessment:

Presentation **last full week of class** (20%)

Submitted **final day of class** (15%)

Longitudinal Participation: (10%)

Course Materials

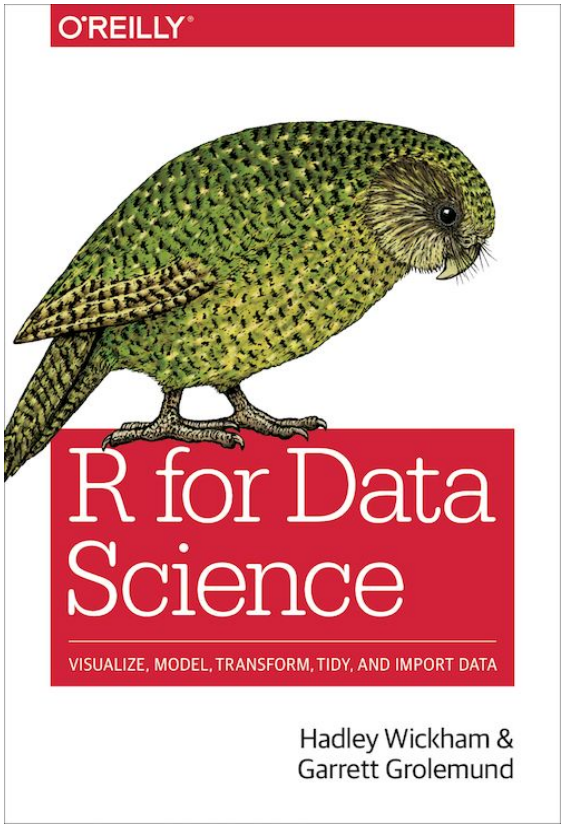

https://r4ds.had.co.nz/

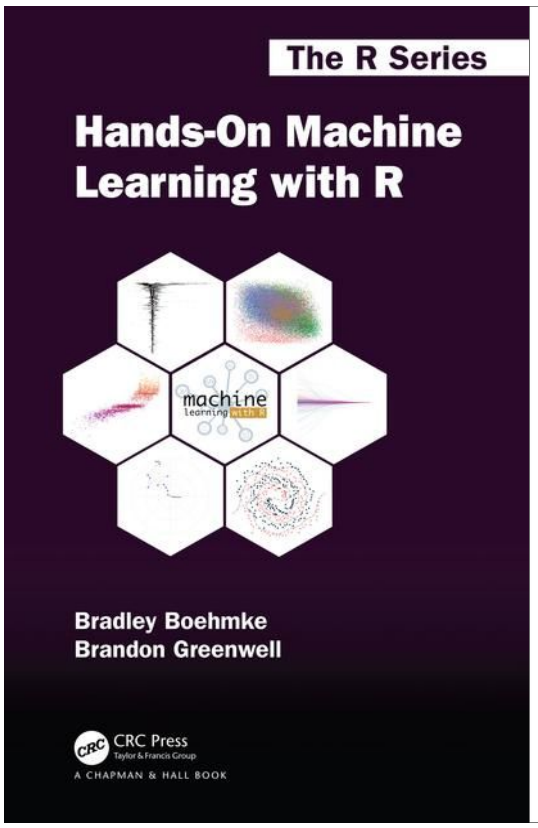

https://bradleyboehmke.github.io/HOML/

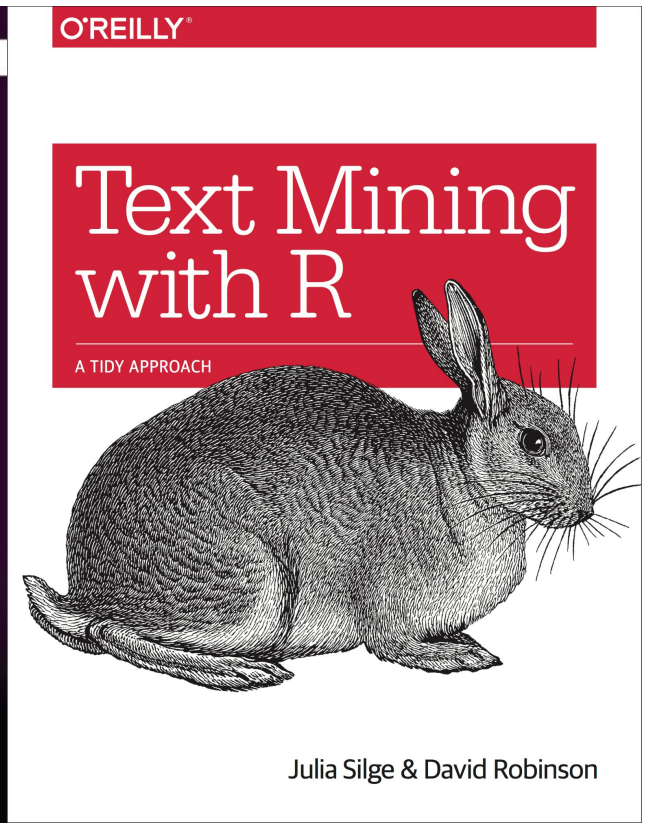

https://www.tidytextmining.com/

Course Website

https://maguire-lab.github.io/health_data_science_research_2025/

Course Website

https://maguire-lab.github.io/health_data_science_research_2025

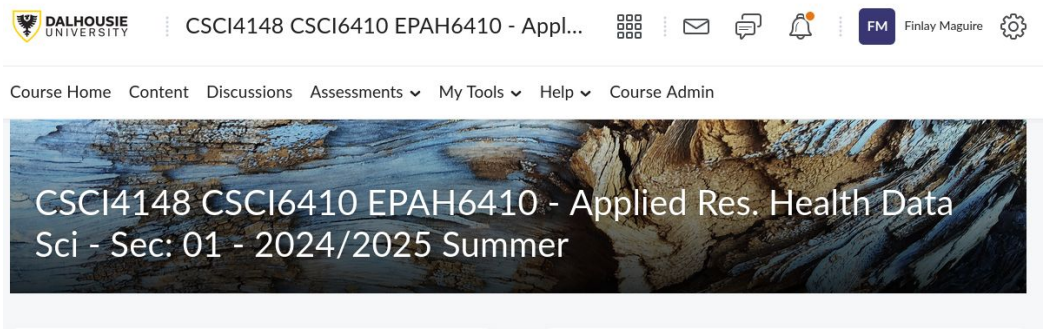

<u>Grades/Submissions:</u>
https://dal.brightspace.com/d2l/home/385844

What is ~~health~~ data science?

Data Science: *Data-intensive interdisciplinary approaches to understand and predict with secondary/live data*

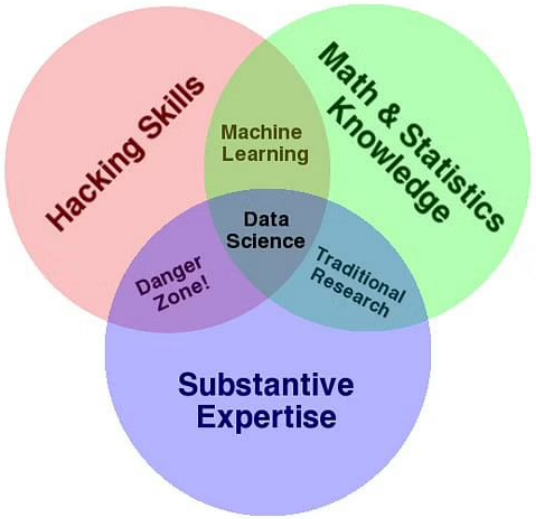

http://drewconway.com/zia/2013/3/26/the-data-science-venn-diagram

Data Science: *Data-intensive interdisciplinary approaches to understand and predict with secondary/live data*

A range of partial and totally overlapping terms:

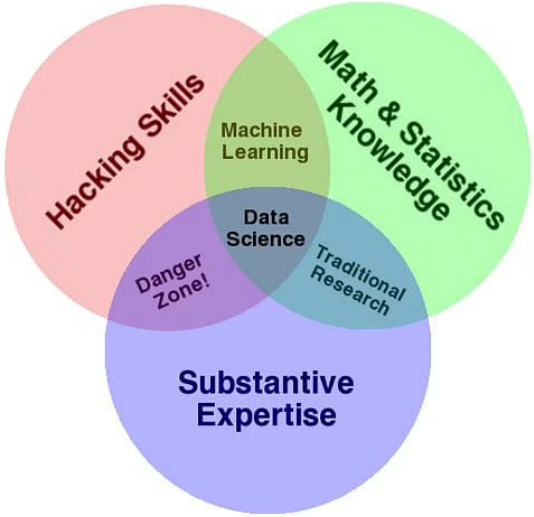

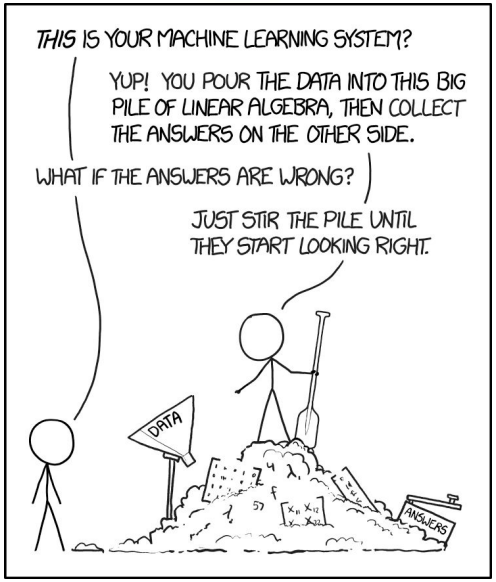

http://drewconway.com/zia/2013/3/26/the-data-science-venn-diagram

Data Science: *Data-intensive interdisciplinary approaches to understand and predict with secondary/live data*

A range of partial and totally overlapping terms:

- Data Analytics
- Data Engineering
- Data Mining
- {Health,Bio,Medical}Informatics
- Database Analysis
- Business Intelligence
- Epidemiology
- Statistics
- **Machine Learning**
- Pattern Recognition
- Predictive Analytics
- Quantitative Researcher
- Scientist
- Analyst
- Algorithmic Modeling

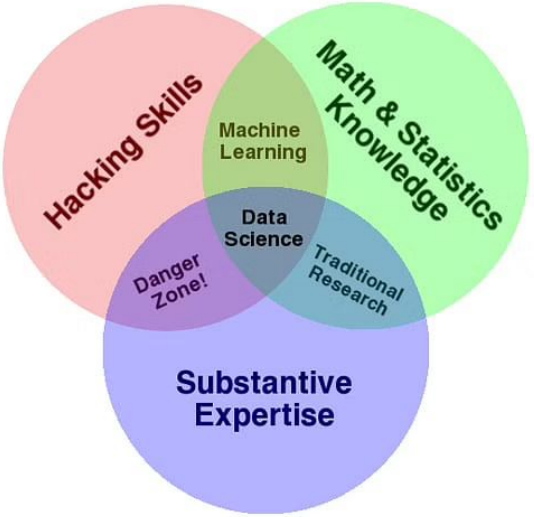

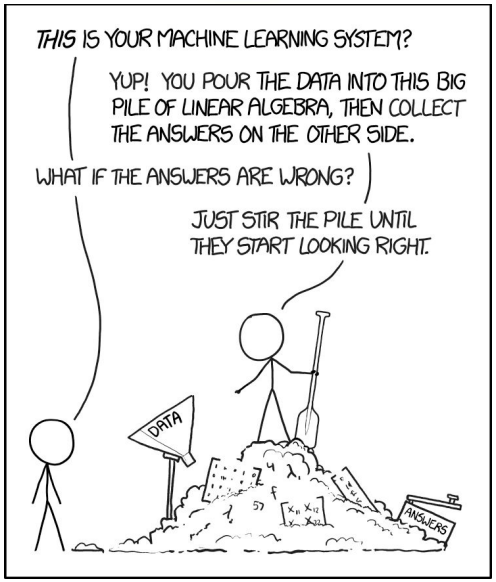

http://drewconway.com/zia/2013/3/26/the-data-science-venn-diagram

So, it is just statistics?

Data Science/ML vs Statistics

- Many shared methods
- Difference in focus/priorities/culture

https://stats.stackexchange.com/questions/442128/machine-learning-vs-statistical-learning-vs-statistics

Data Science/ML vs Statistics

- Many shared methods
- Difference in focus/priorities/culture
- Statistics ~ tries to understand how outcome was generated by data
- ML infers/learns a process for linking data to outcome
- Alternative framing:
  - Data Modelling vs Algorithmic Modelling

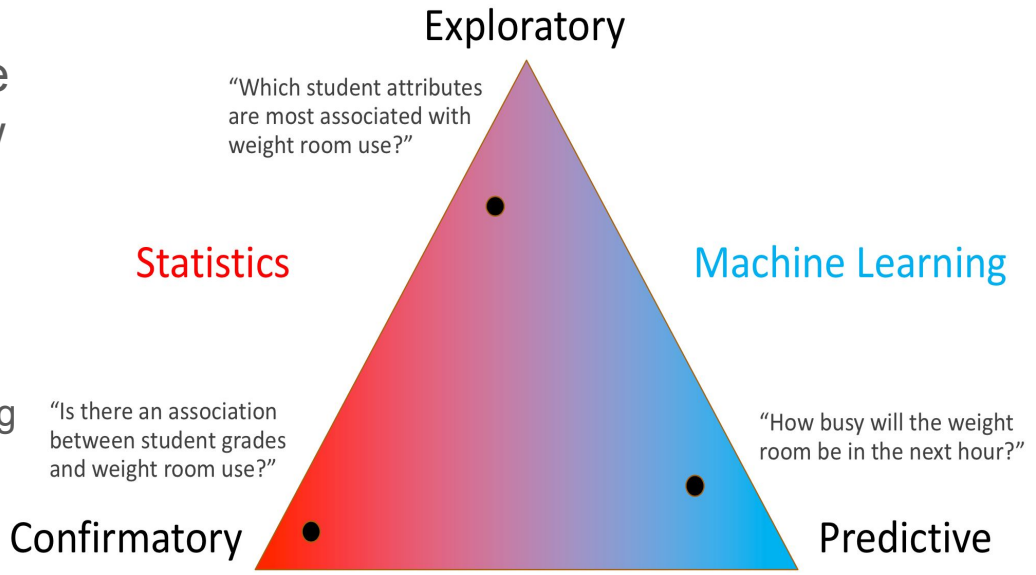

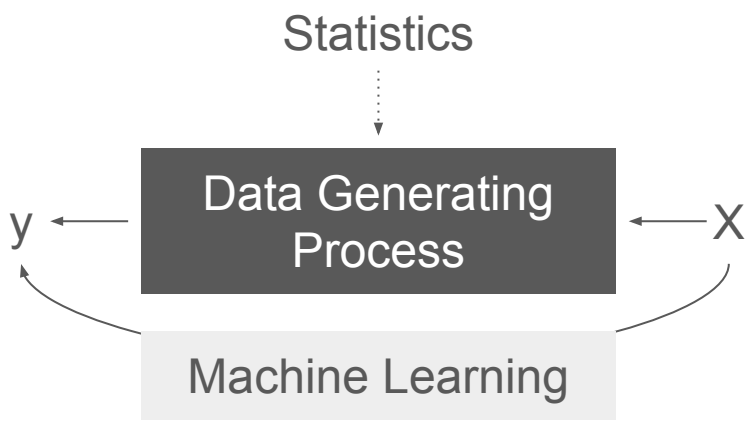

https://stats.stackexchange.com/questions/442128/machine-learning-vs-statistical-learning-vs-statistics

Data Science/ML vs Statistics

- Many shared methods
- Difference in focus/priorities/culture
- Statistics ~ tries to understand how outcome was generated by data
- ML infers/learns a process for linking data to outcome
- Alternative framing:
  - Data Modelling vs Algorithmic Modelling

- DS/ML Pitfalls (can be):
  - Less rigorous/principled
  - Prone to reinventing the wheel
- DS/ML Benefits (can be):
  - More flexible
  - Less prescriptive/intimidating

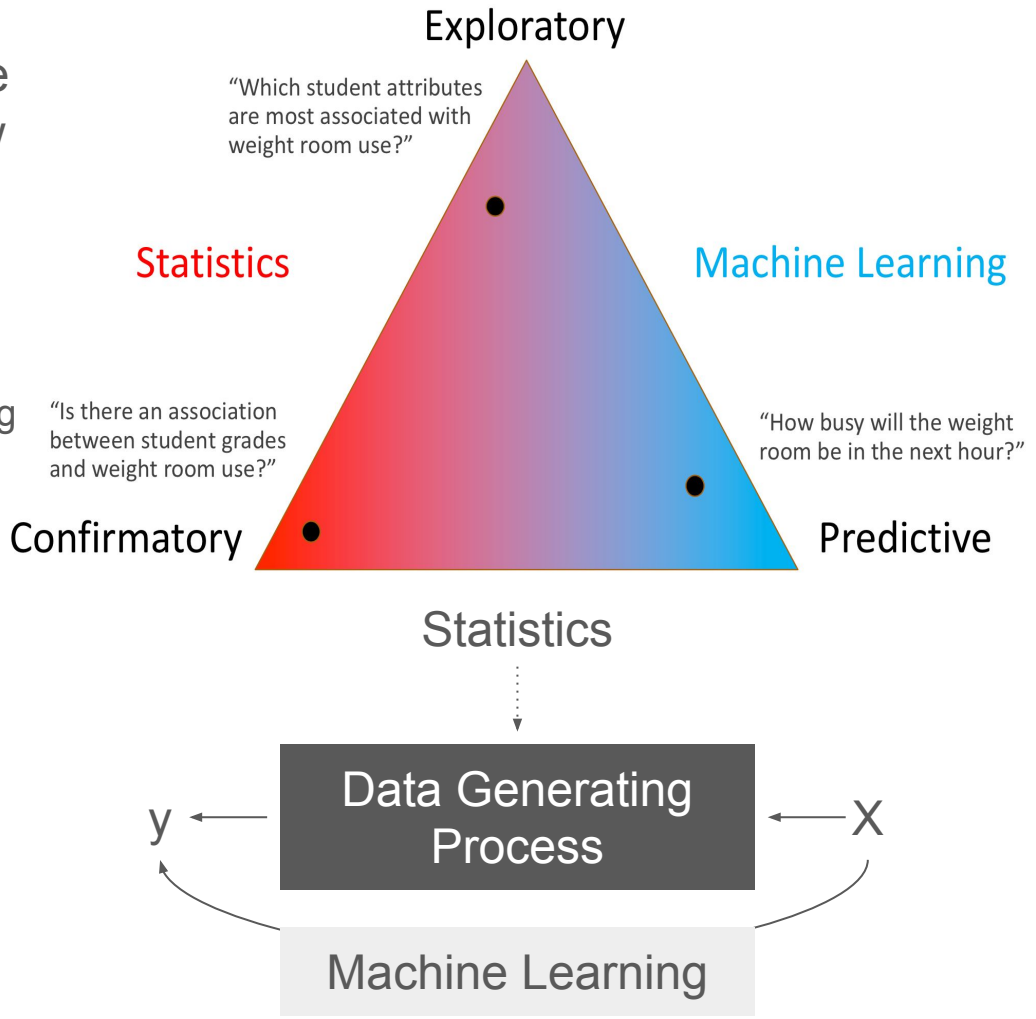

https://stats.stackexchange.com/questions/442128/machine-learning-vs-statistical-learning-vs-statistics

Data science centers exploratory data analysis

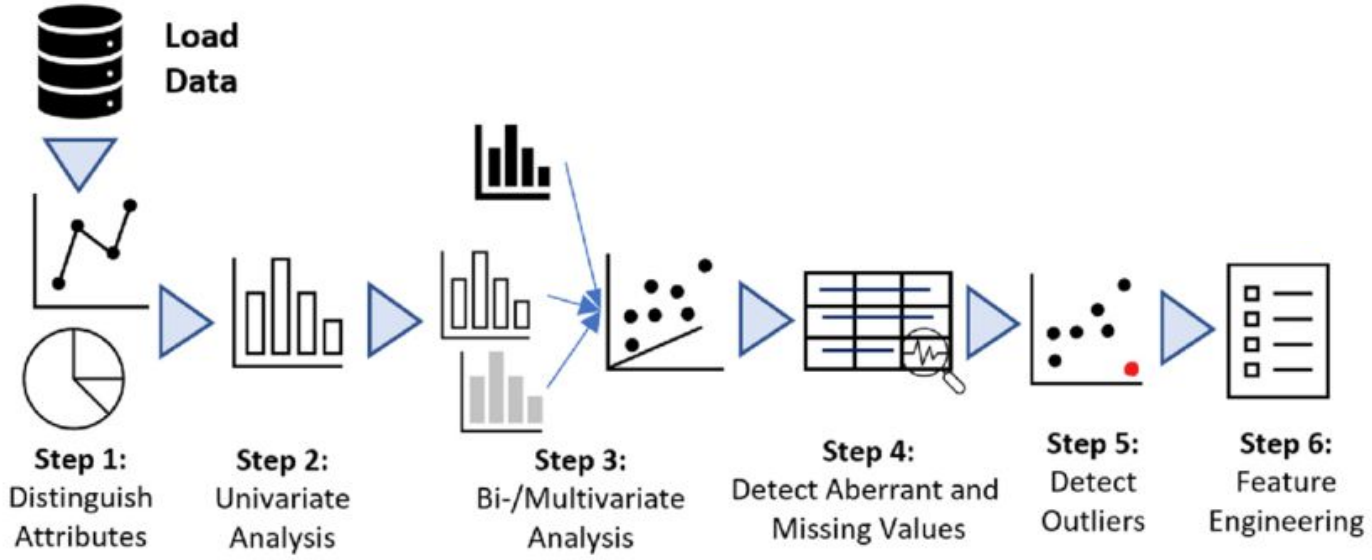

10.3390/su12124995

Data science supports inductive approaches

Data science supports inductive approaches

Deductive:

- “Condition X, causes Y”
- Collect data
- Perform (typically) frequentist statistical tests
- Reject or confirm null hypothesis

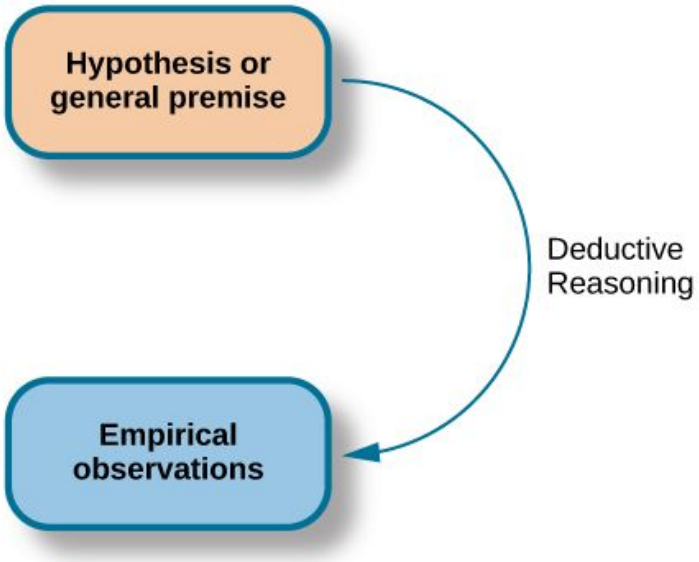

https://opened.cuny.edu/courseware/lesson/14/student/?task=3

Data science supports inductive approaches

Deductive:

- “Condition X, causes Y”
- Collect data
- Perform (typically) frequentist statistical tests
- Reject or confirm null hypothesis

Inductive:

- Collect data
- Identify patterns in the data
- Observe X and Y seem connected somehow
- Quantify strength of association e.g., prediction performance

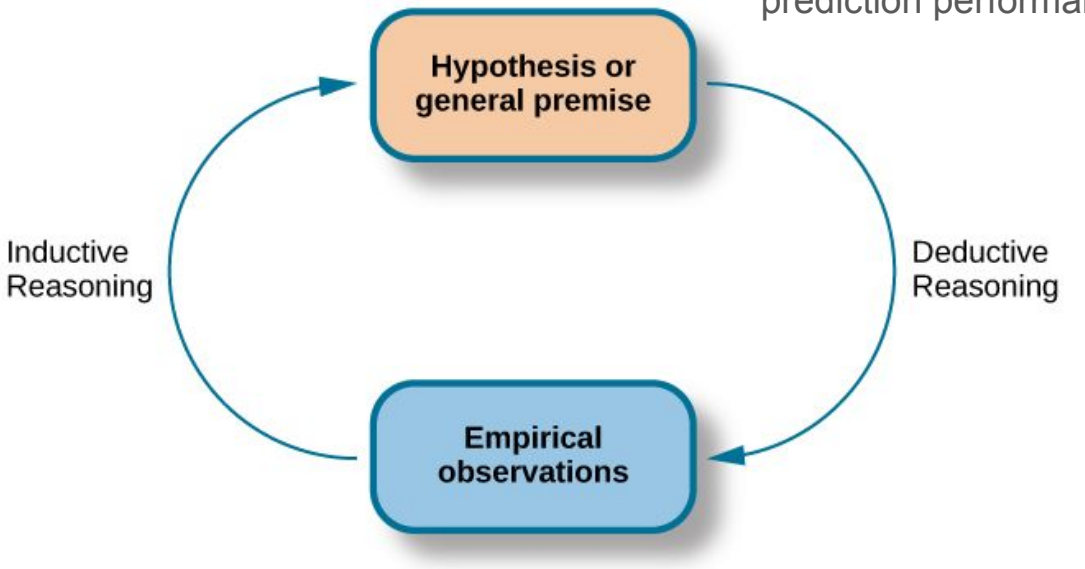

https://opened.cuny.edu/courseware/lesson/14/student/?task=3

Data science aligns with knowledge cycle

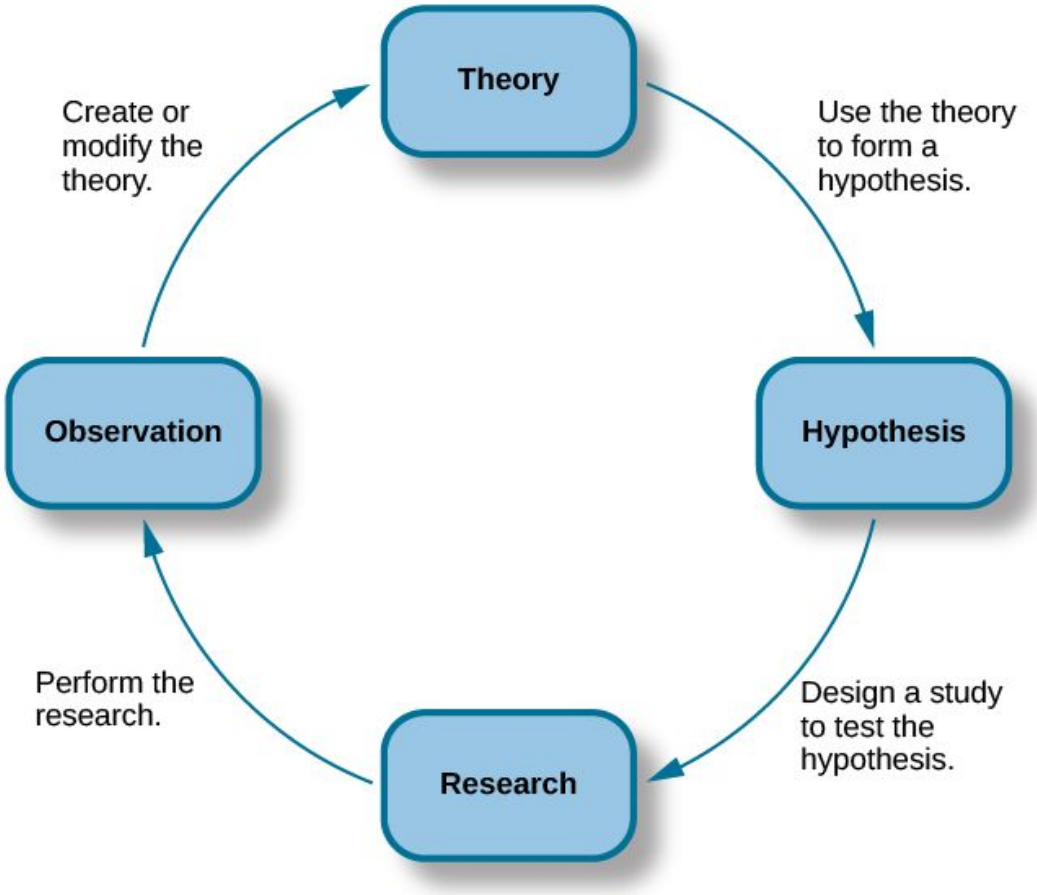

https://opened.cuny.edu/courseware/lesson/14/student/?task=3

Data science aligns with knowledge cycle

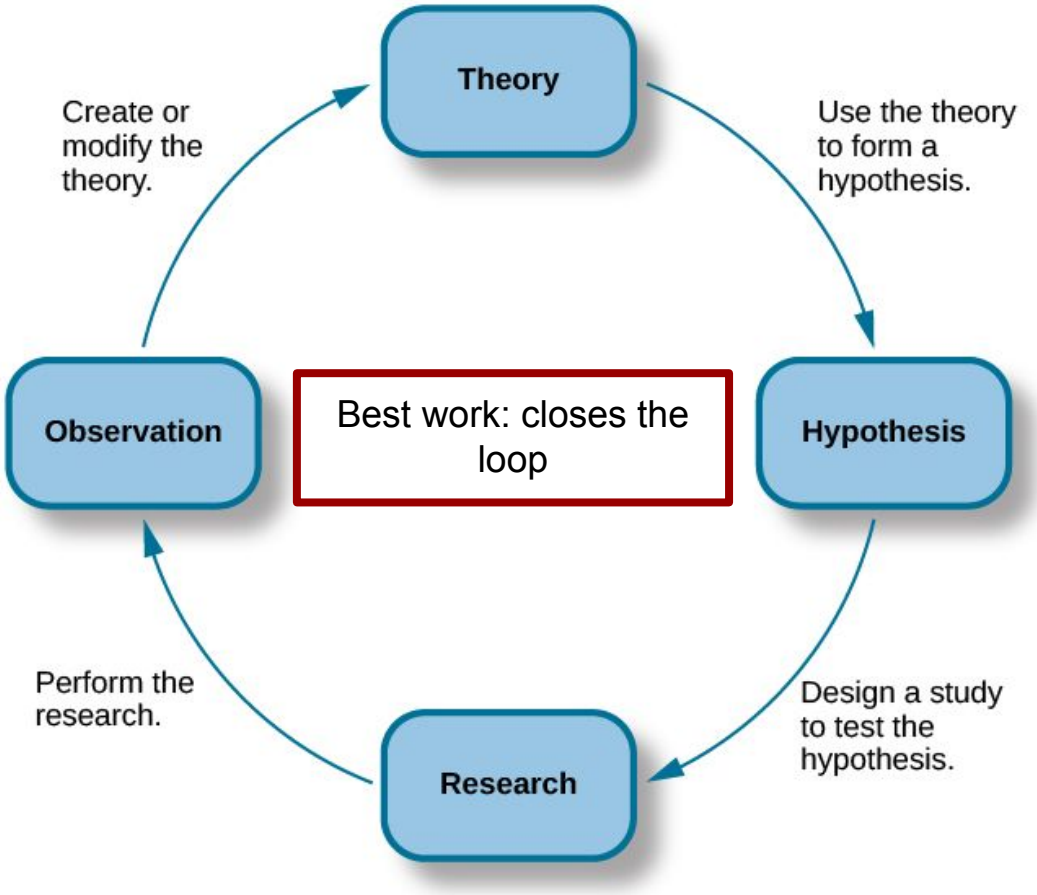

https://opened.cuny.edu/courseware/lesson/14/student/?task=3

Data science is integrated into a data ecosystem

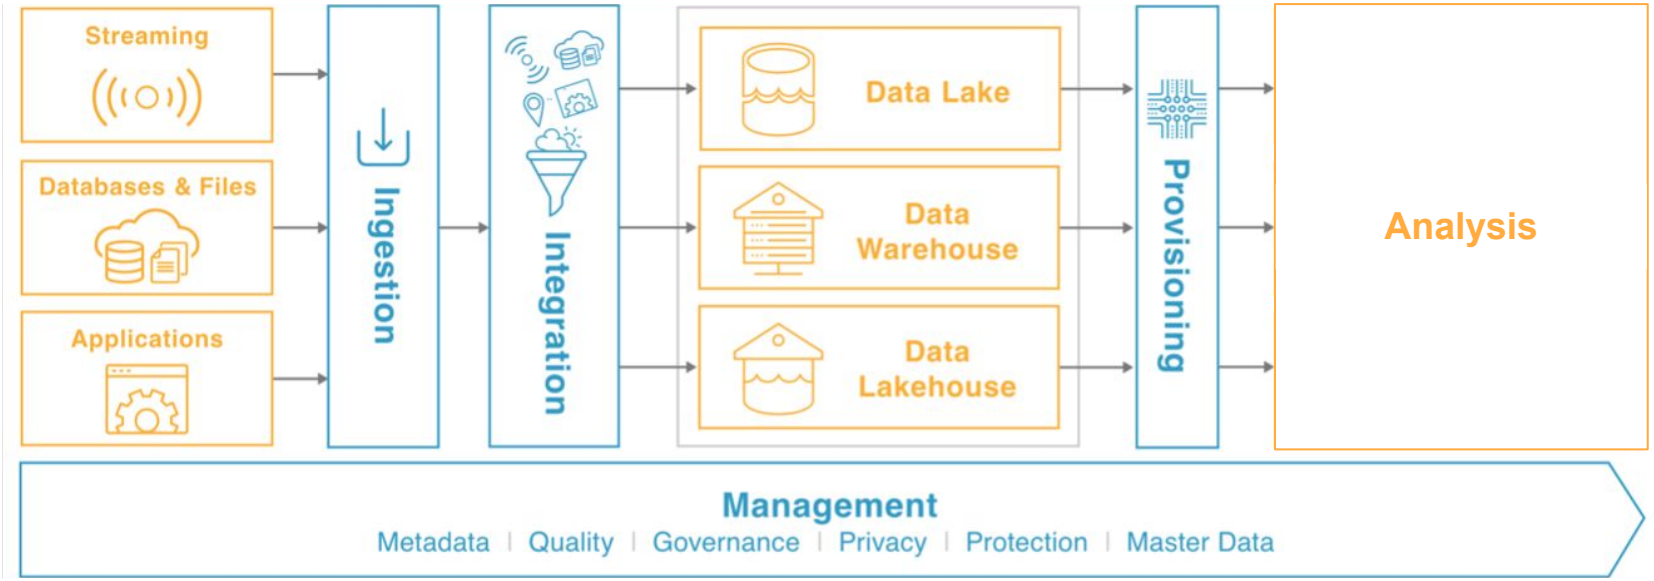

https://www.2ndwatch.com/blog/what-is-a-data-pipeline-and-how-to-build-one/

Data science is integrated into a data ecosystem

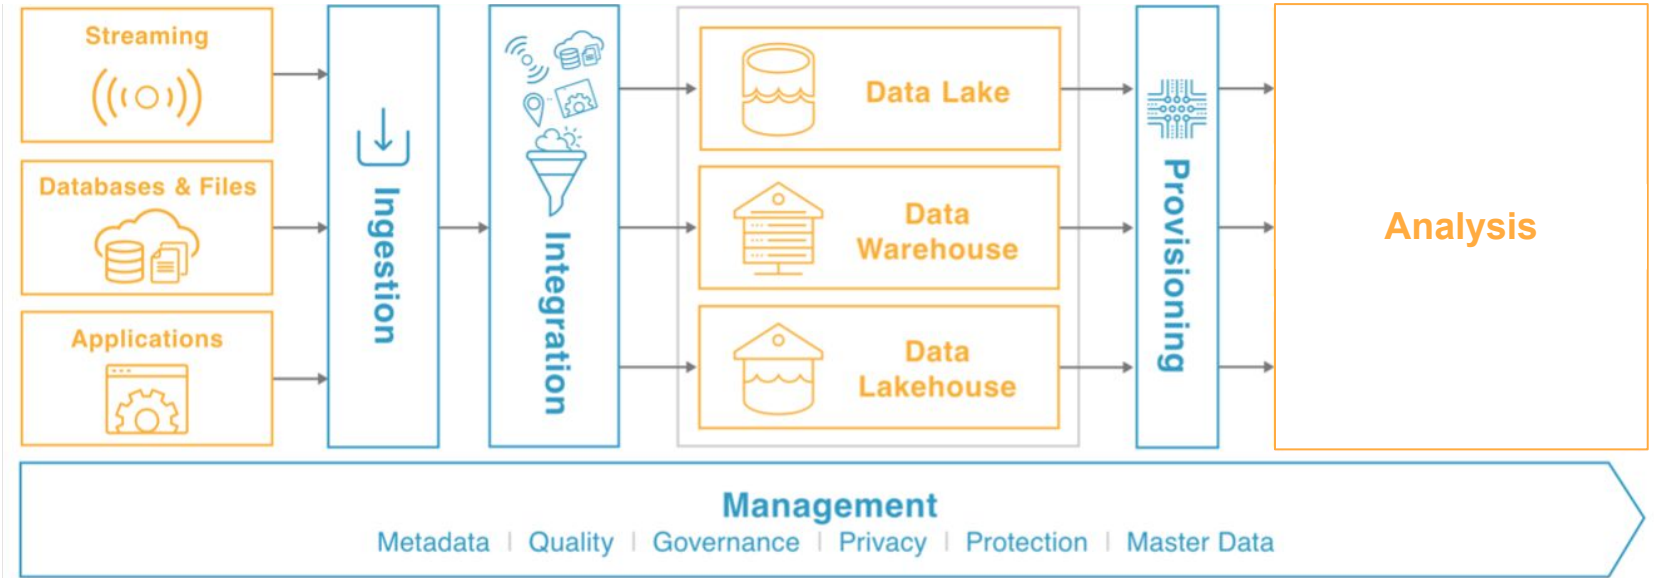

https://www.2ndwatch.com/blog/what-is-a-data-pipeline-and-how-to-build-one/

Some Open-Source Orchestration Tools:

Apache Airflow, DAGSTER, DVC, Flyte, Kale, Kedro, Kubeflow, Luigi, METAFLOW, Ploomber, TensorFlow Extended

https://ploomber.io/blog/survey/

OK, what is **Health** Data Science?

Data Science applied to Health Data

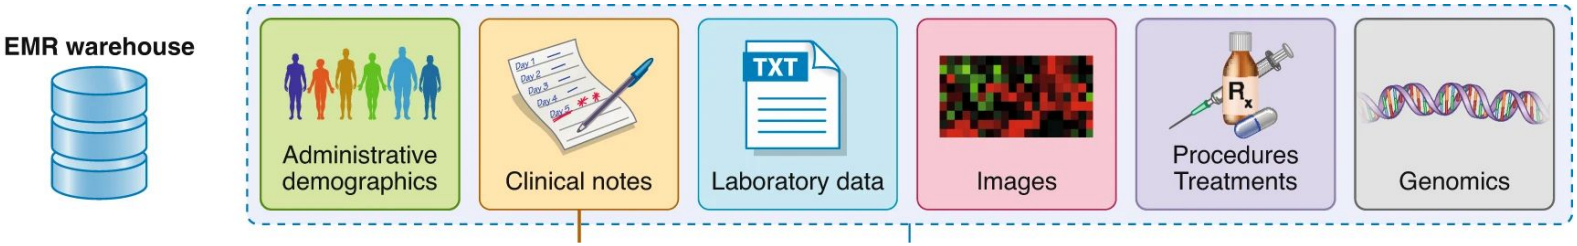

Why “health data” instead of “medical data”: health encompasses medical ***(contentious)***

Data Science applied to Health Data

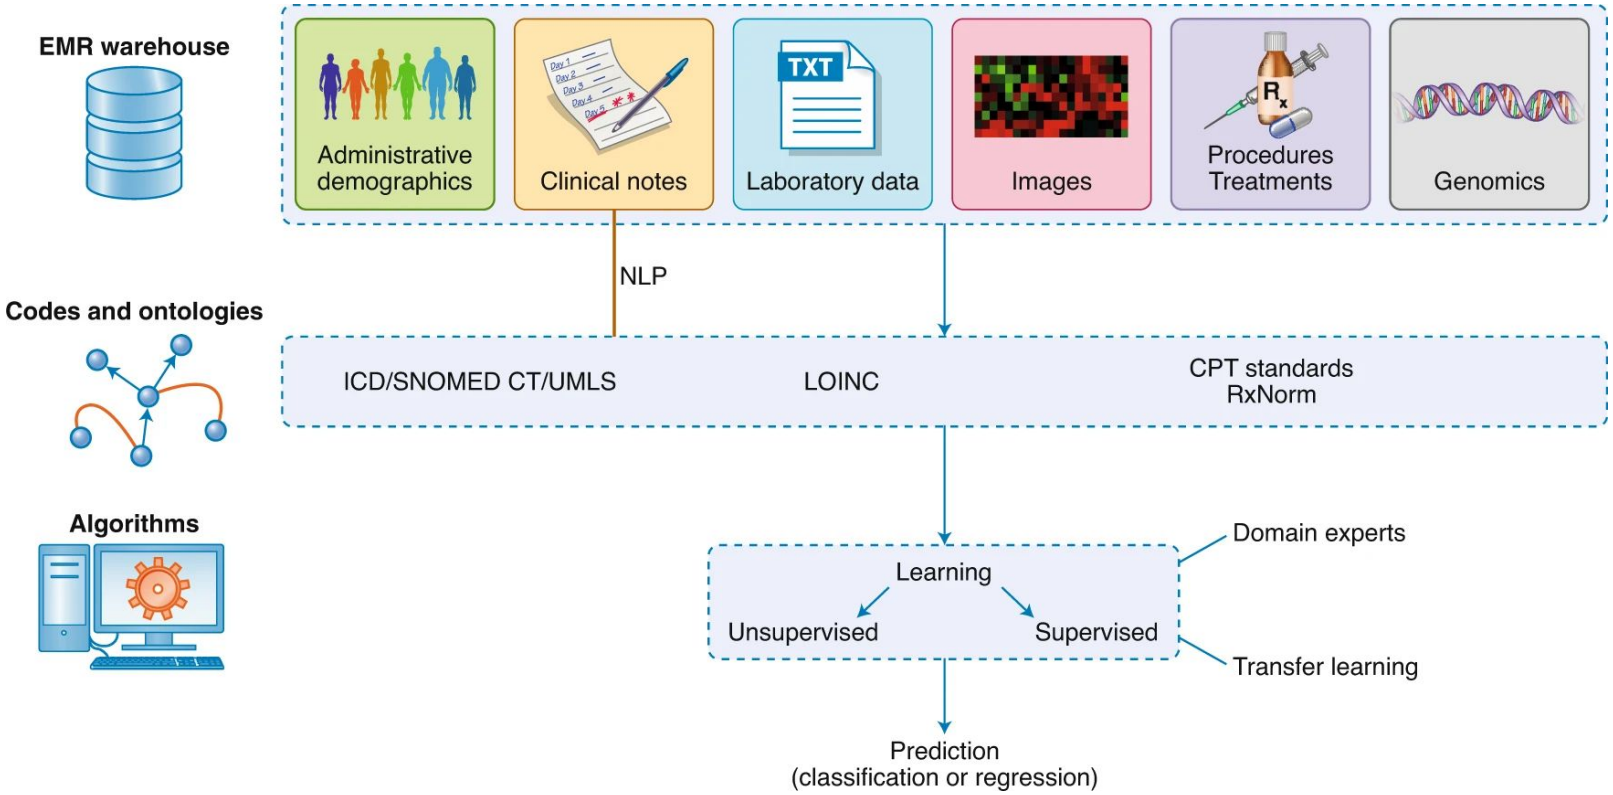

https://www.nature.com/articles/s41588-020-0698-y/figures/2

Why “health data” instead of “medical data”: health encompasses medical ***(contentious)***

Opportunity of Health Data Science

Benefits (and pitfalls!) of data science in general combined with:

- Huge amounts of health data

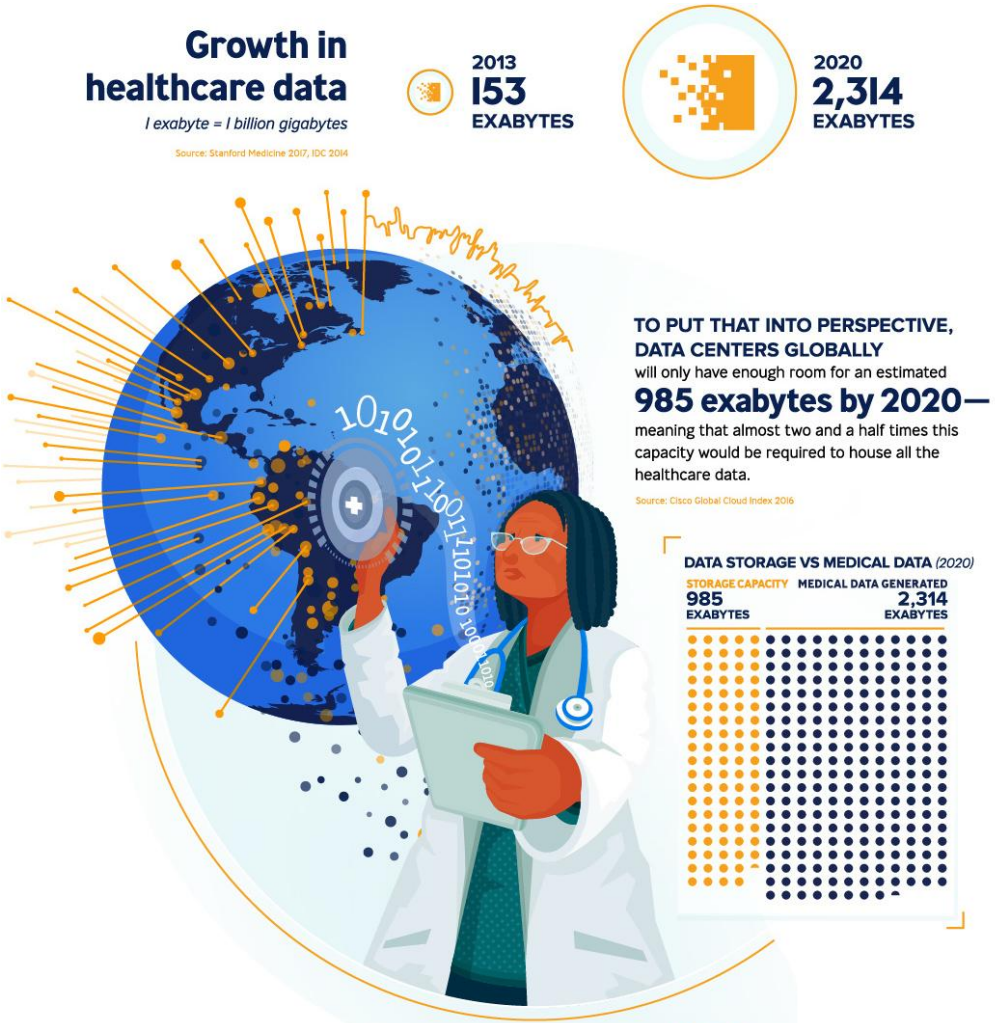

https://www.visualcapitalist.com/big-data-healthcare/

Opportunity of Health Data Science

Benefits (and pitfalls!) of data science in general combined with:

- Huge amounts of health data
- Many **interesting** and **important problems**

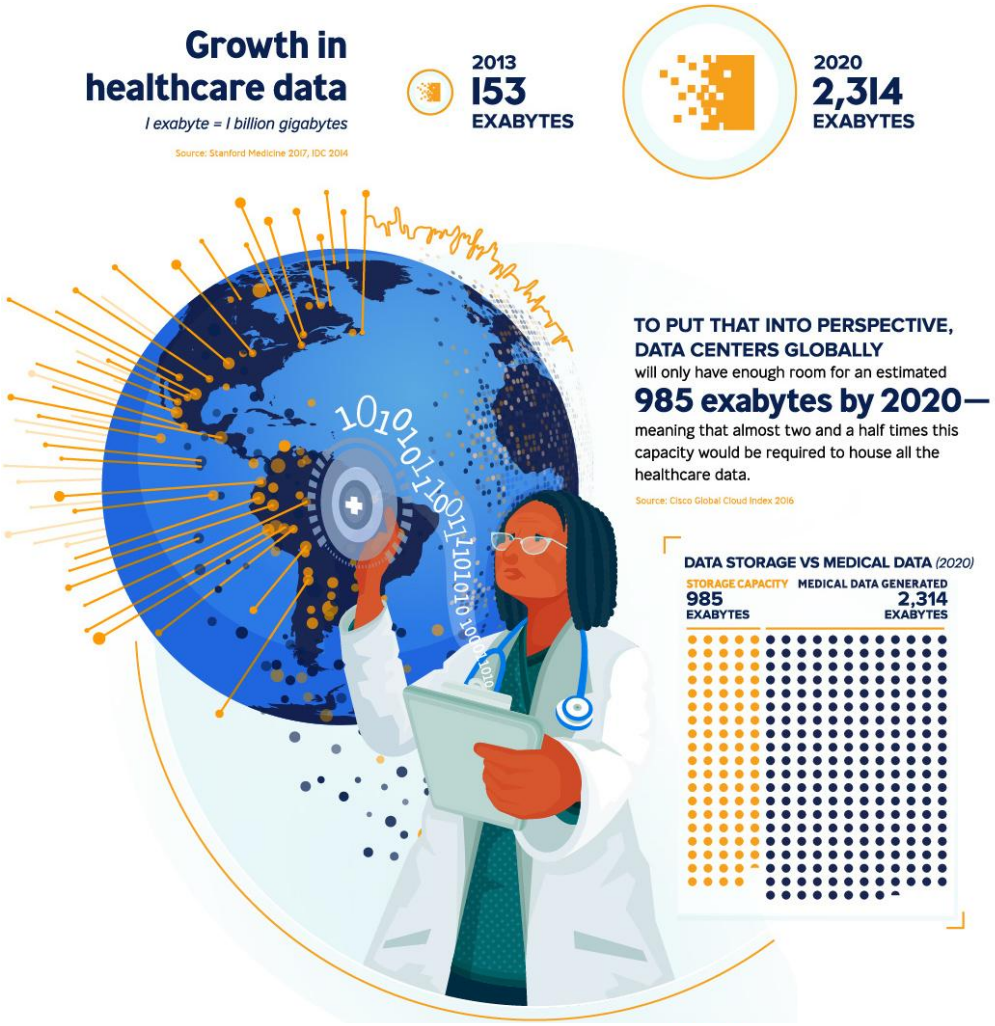

https://www.visualcapitalist.com/big-data-healthcare/

Opportunity of Health Data Science

Benefits (and pitfalls!) of data science in general combined with:

- Huge amounts of health data
- Many **interesting** and **important problems**
- Many domain experts desperate for data-related help with these problems

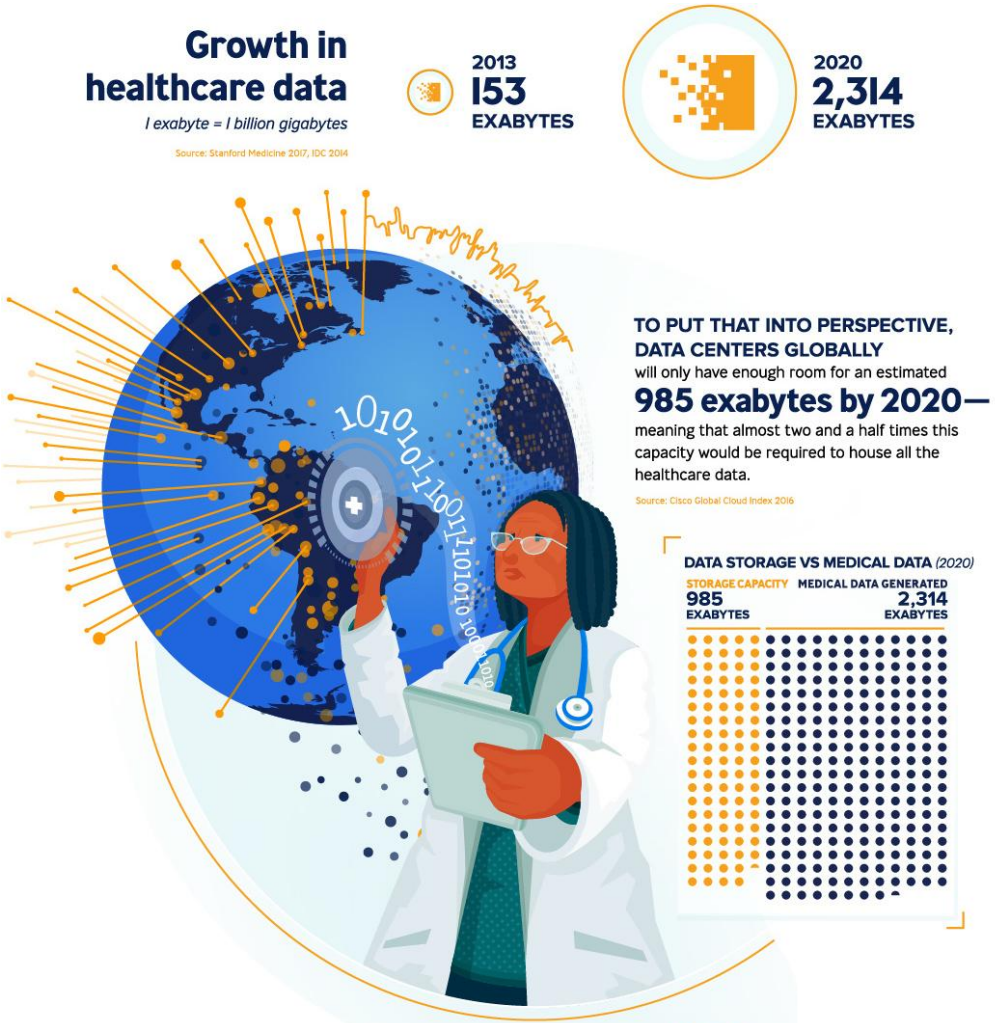

https://www.visualcapitalist.com/big-data-healthcare/

Opportunity of Health Data Science

Benefits (and pitfalls!) of data science in general combined with:

- Huge amounts of health data
- Many **interesting** and **important problems**
- Many domain experts desperate for data-related help with these problems
- Relative few skilled data science practitioners

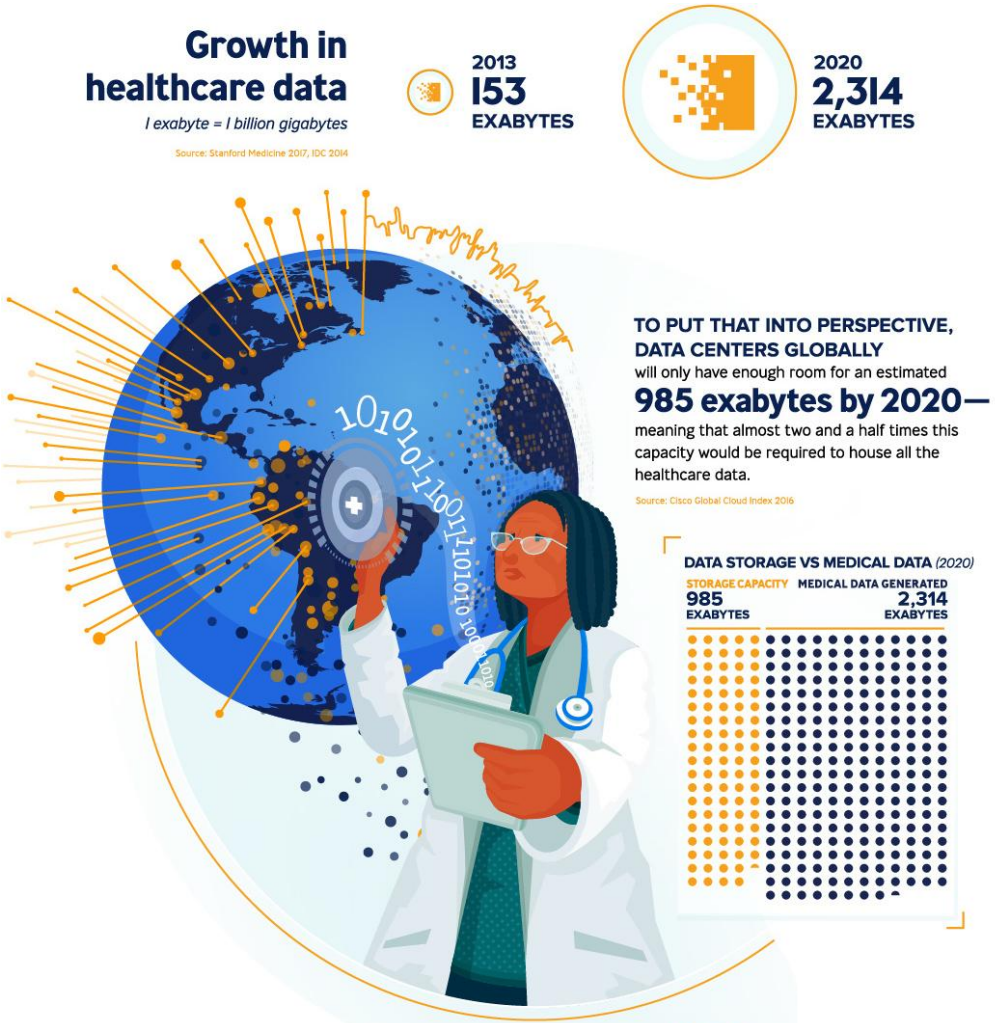

https://www.visualcapitalist.com/big-data-healthcare/

(Some) Challenges of Health Data Science

- Lots of hype

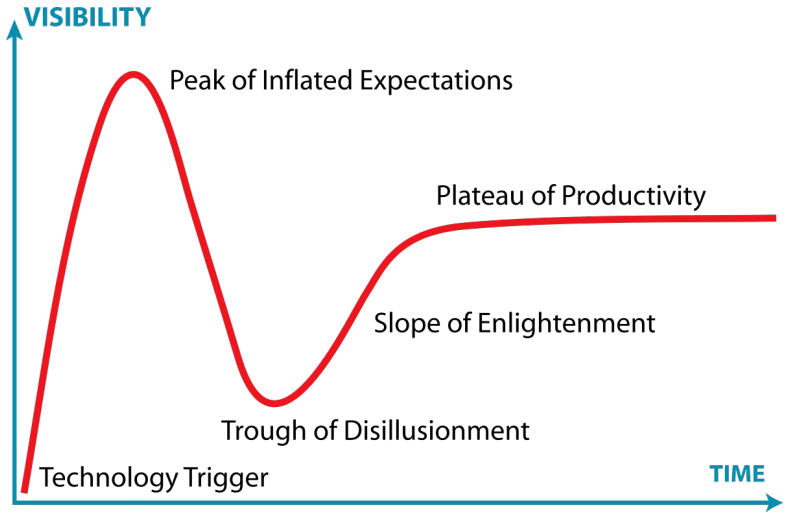

(Some) Challenges of Health Data Science

- Lots of hype
- Lots of grifters

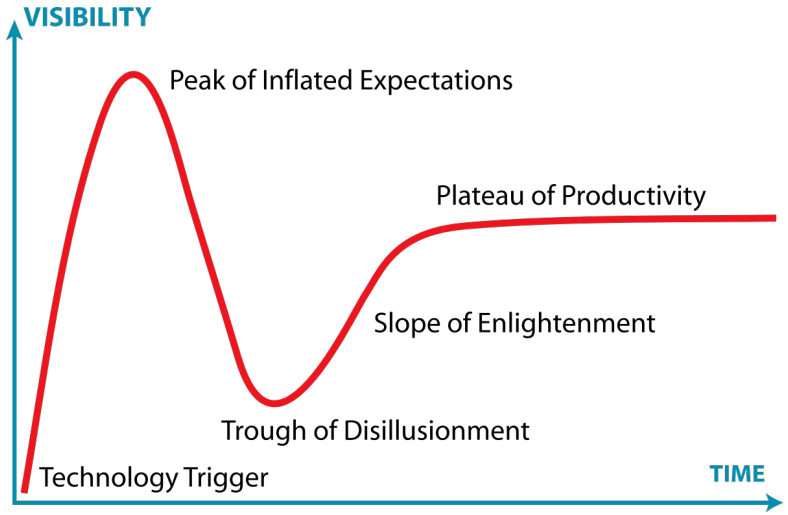

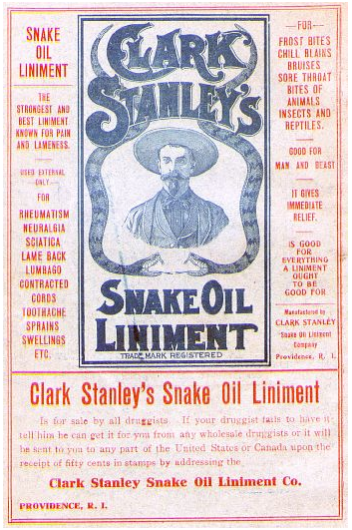

https://upload.wikimedia.org/wikipedia/commons/9/94/Gartner_Hype_Cycle.svg
https://commons.wikimedia.org/wiki/File:Clark_Stanley%27s_Snake_Oil_Liniment.png

(Some) Challenges of Health Data Science

- Lots of hype
- Lots of grifters
- Data quality issues
- Contextual/Metadata quality issues

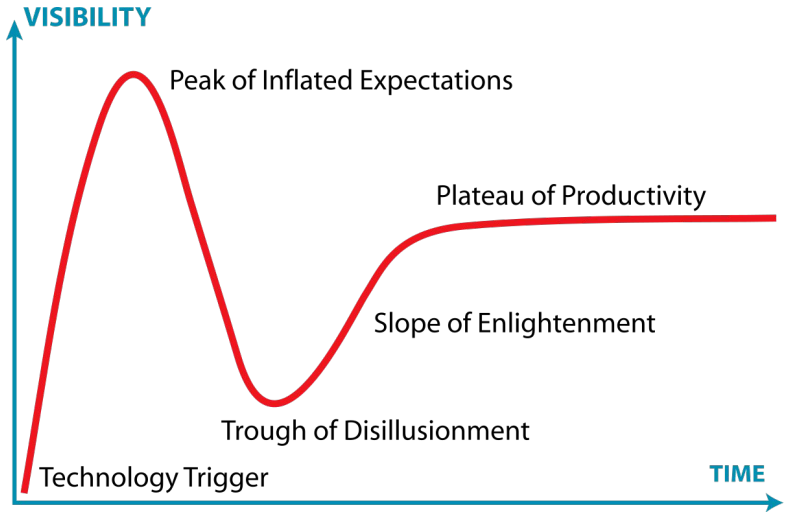

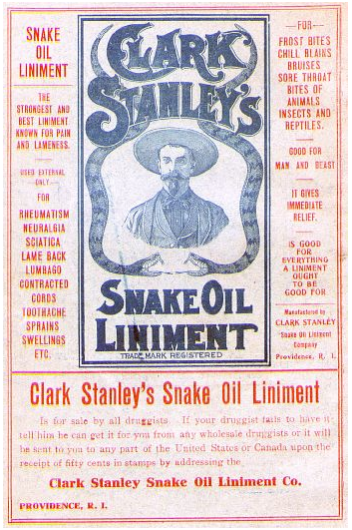

https://upload.wikimedia.org/wikipedia/commons/9/94/Gartner_Hype_Cycle.svg
https://commons.wikimedia.org/wiki/File:Clark_Stanley%27s_Snake_Oil_Liniment.png

(Some) Challenges of Health Data Science

- Lots of hype
- Lots of grifters
- Data quality issues
- Contextual/Metadata quality issues
- Regulatory challenges
- Influence of US health system
- Ethical pitfalls
- Treatment to the mean
- Knowledge Translation and Operations: **Hard**

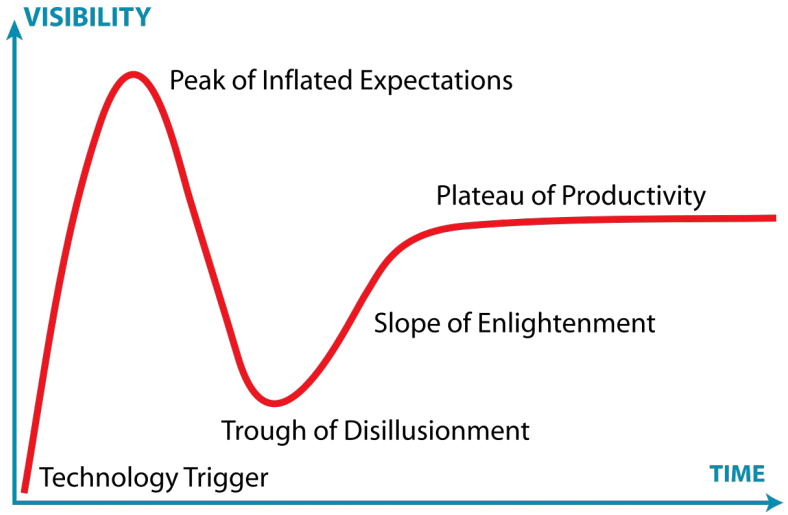

https://www.r-bloggers.com/2019/08/new-course-learn-advanced-data-cleaning-in-r/

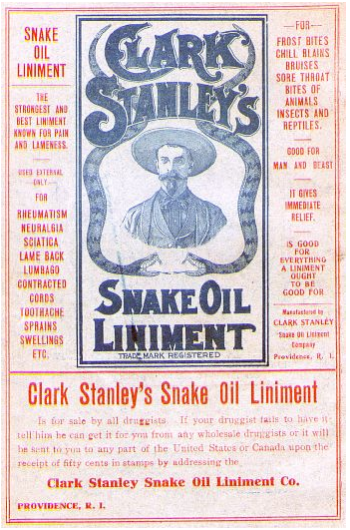

https://upload.wikimedia.org/wikipedia/commons/9/94/Gartner_Hype_Cycle.svg
https://commons.wikimedia.org/wiki/File:Clark_Stanley%27s_Snake_Oil_Liniment.png

What parts of health data science will this course cover?

What parts of health data science will this course cover?

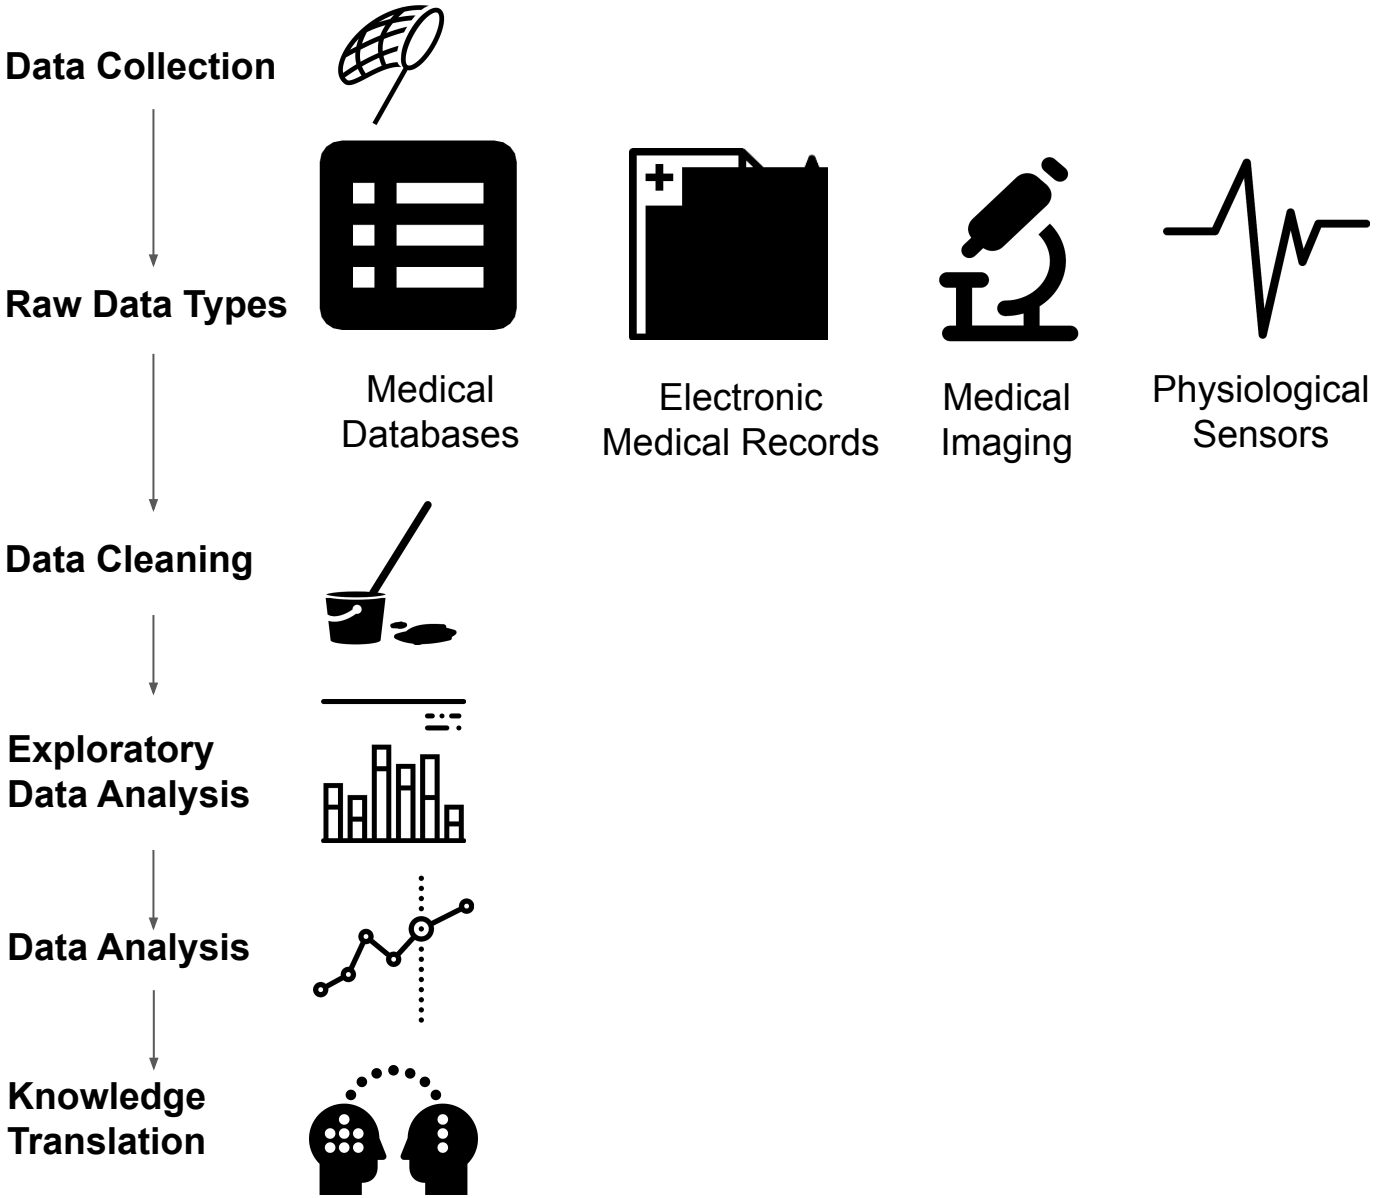

What parts of health data science will this course cover?

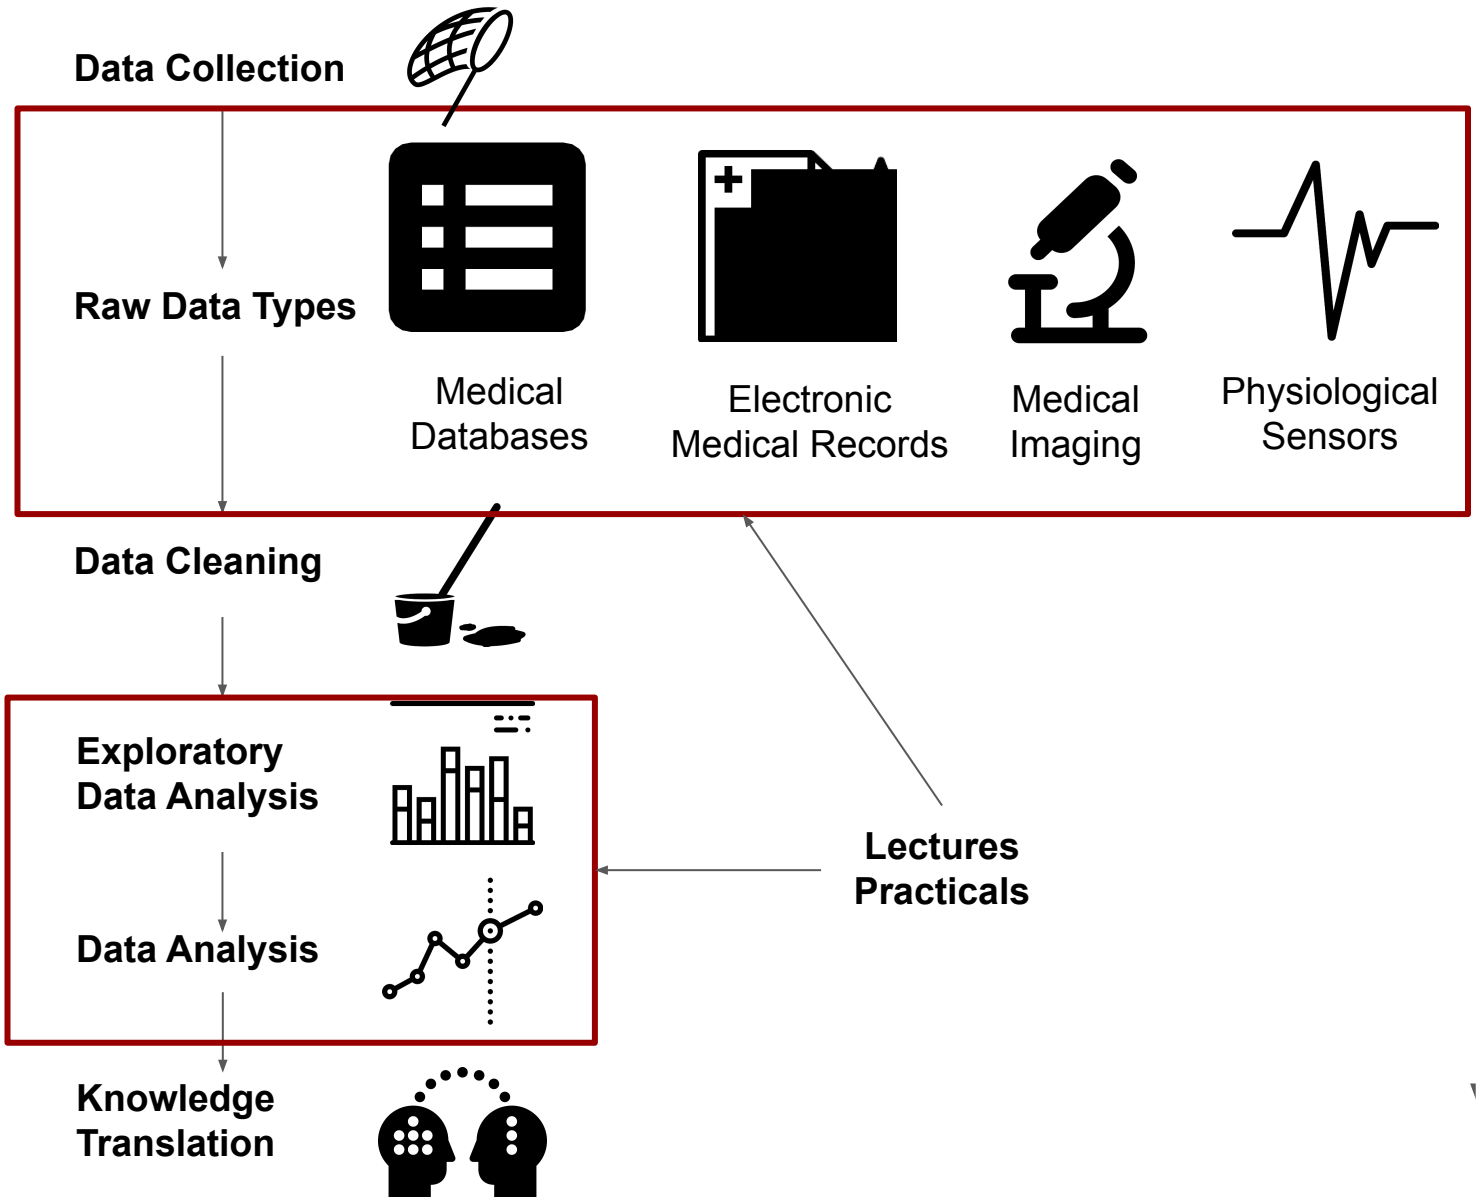

What parts of health data science will this course cover?

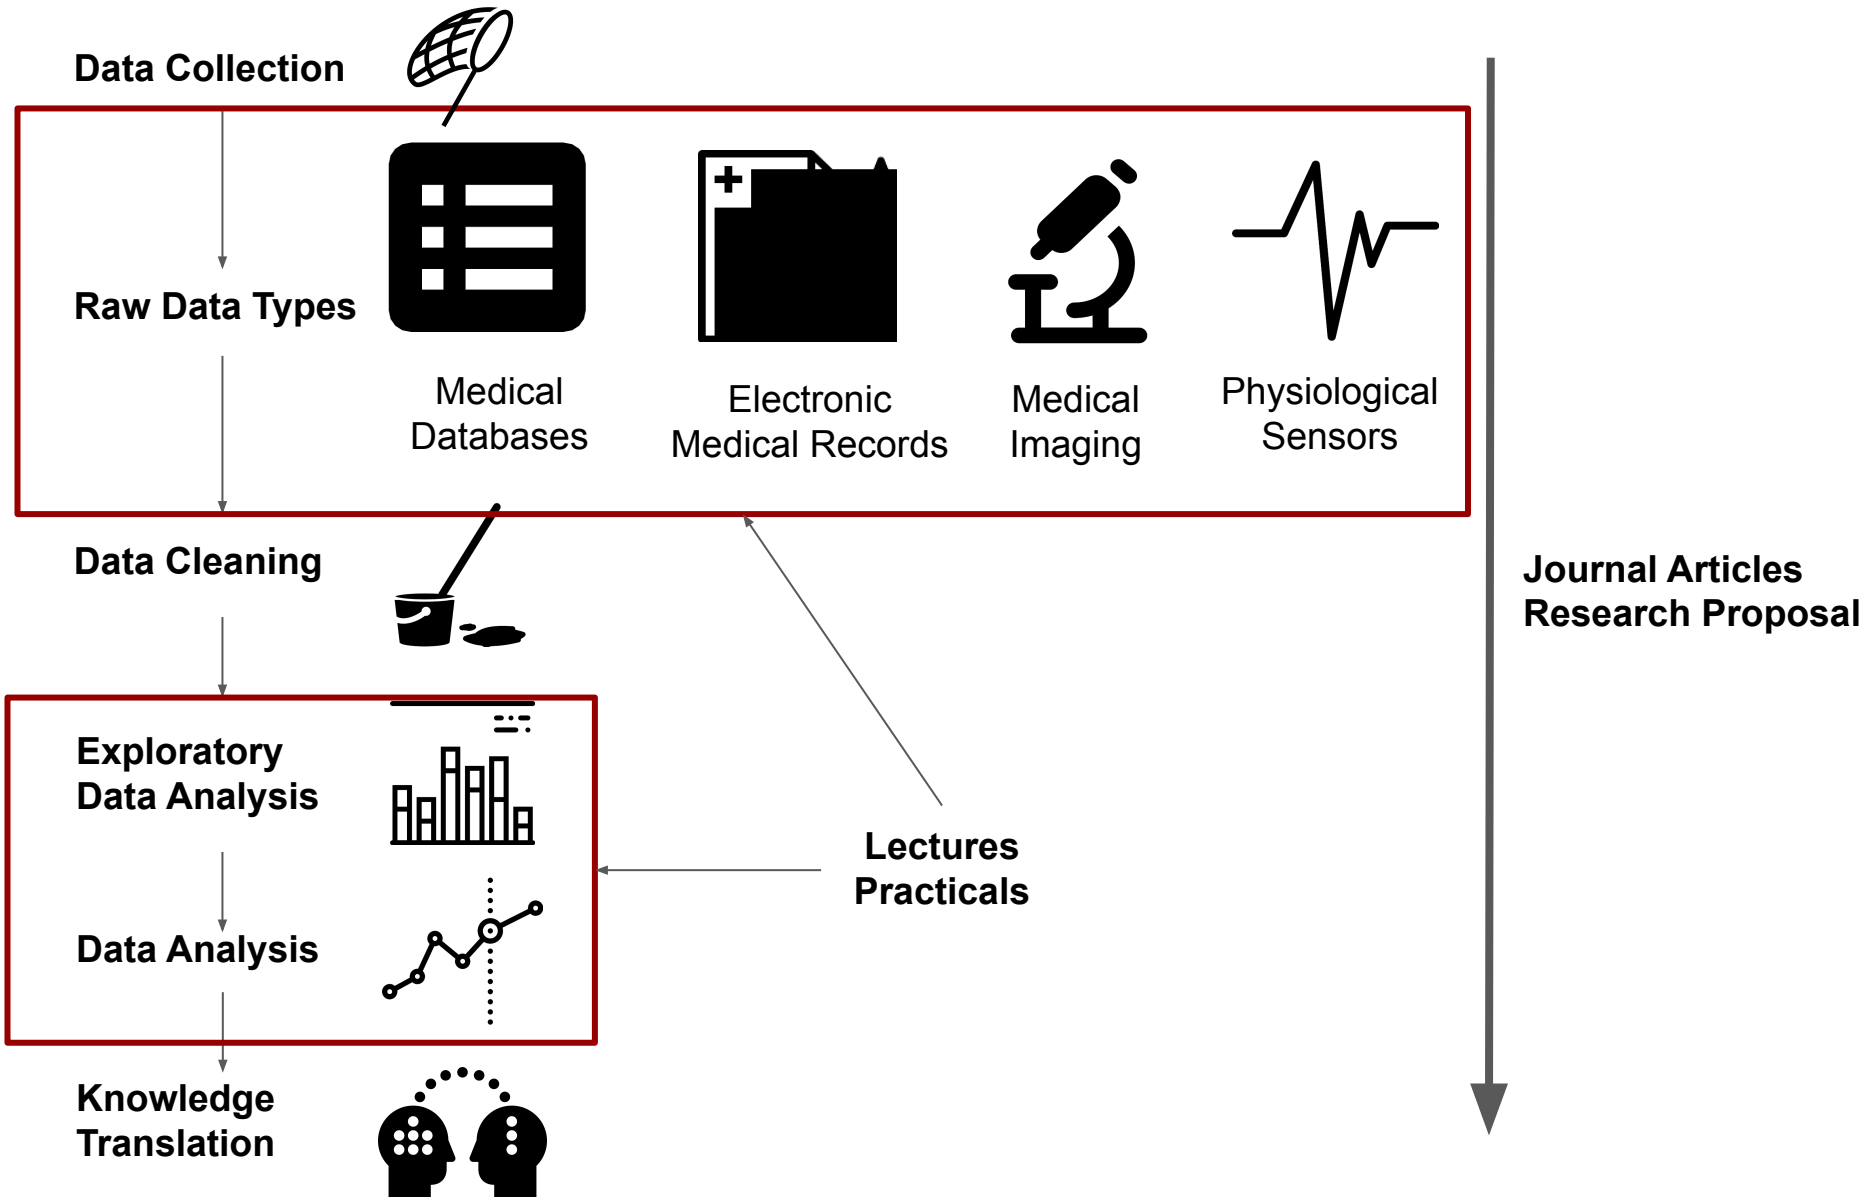

Let’s take a 5 minute break!

Tools for Reproducible Health Data Science

Rstudio, Rmarkdown, Git

Why do we care about reproducibility?

Reproducibility should be the bare minimum

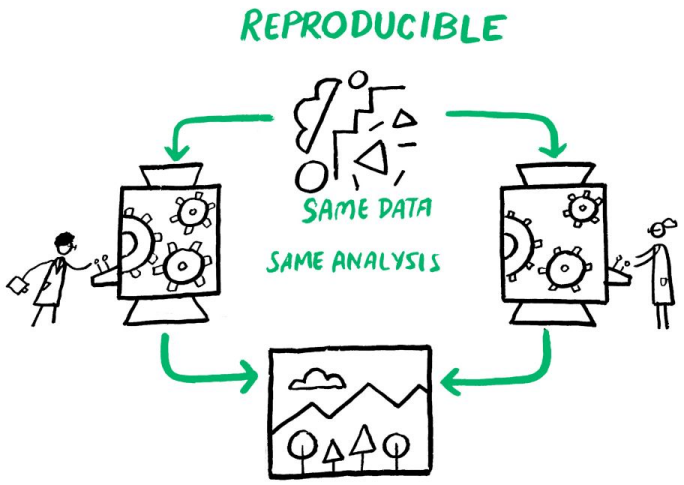

Reproducibility should be the bare minimum

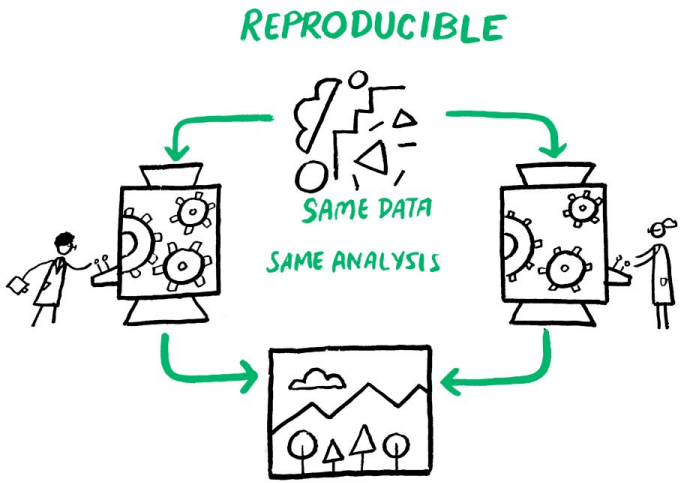

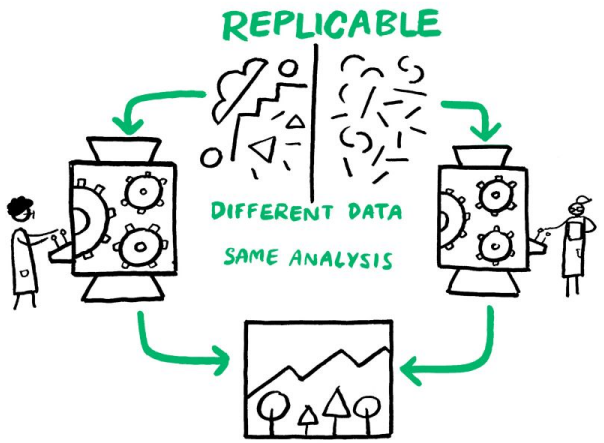

Reproducibility should be the bare minimum

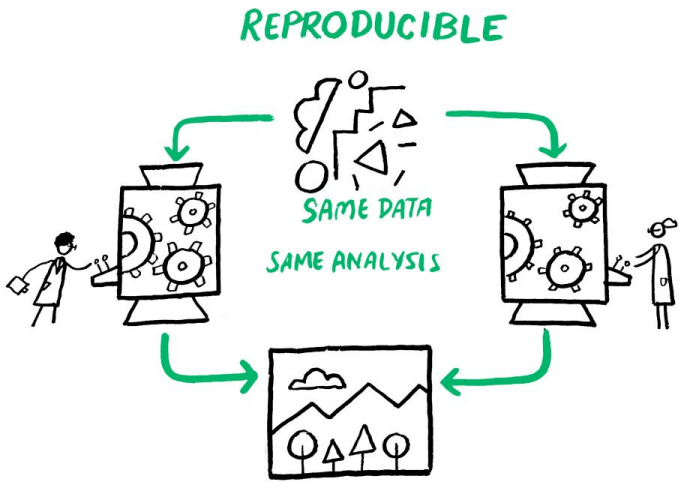

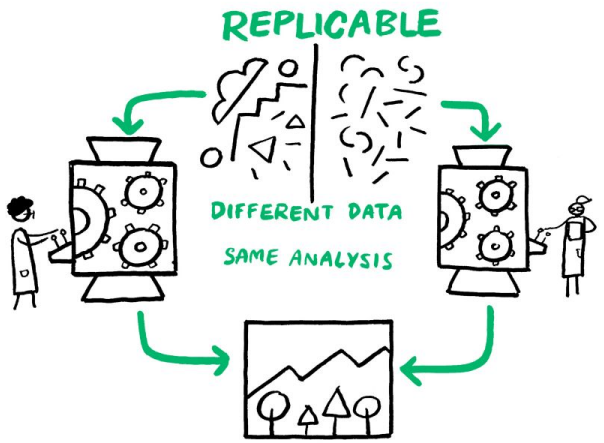

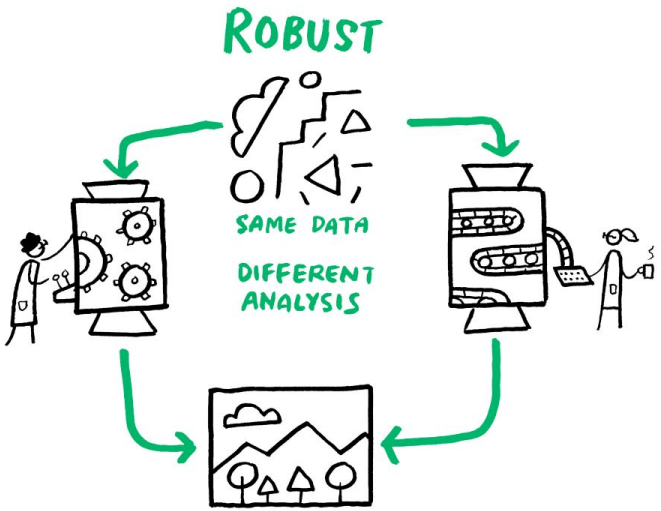

oliviergimenez.github.io/reproducible-science-workshop

Reproducibility should be the bare minimum

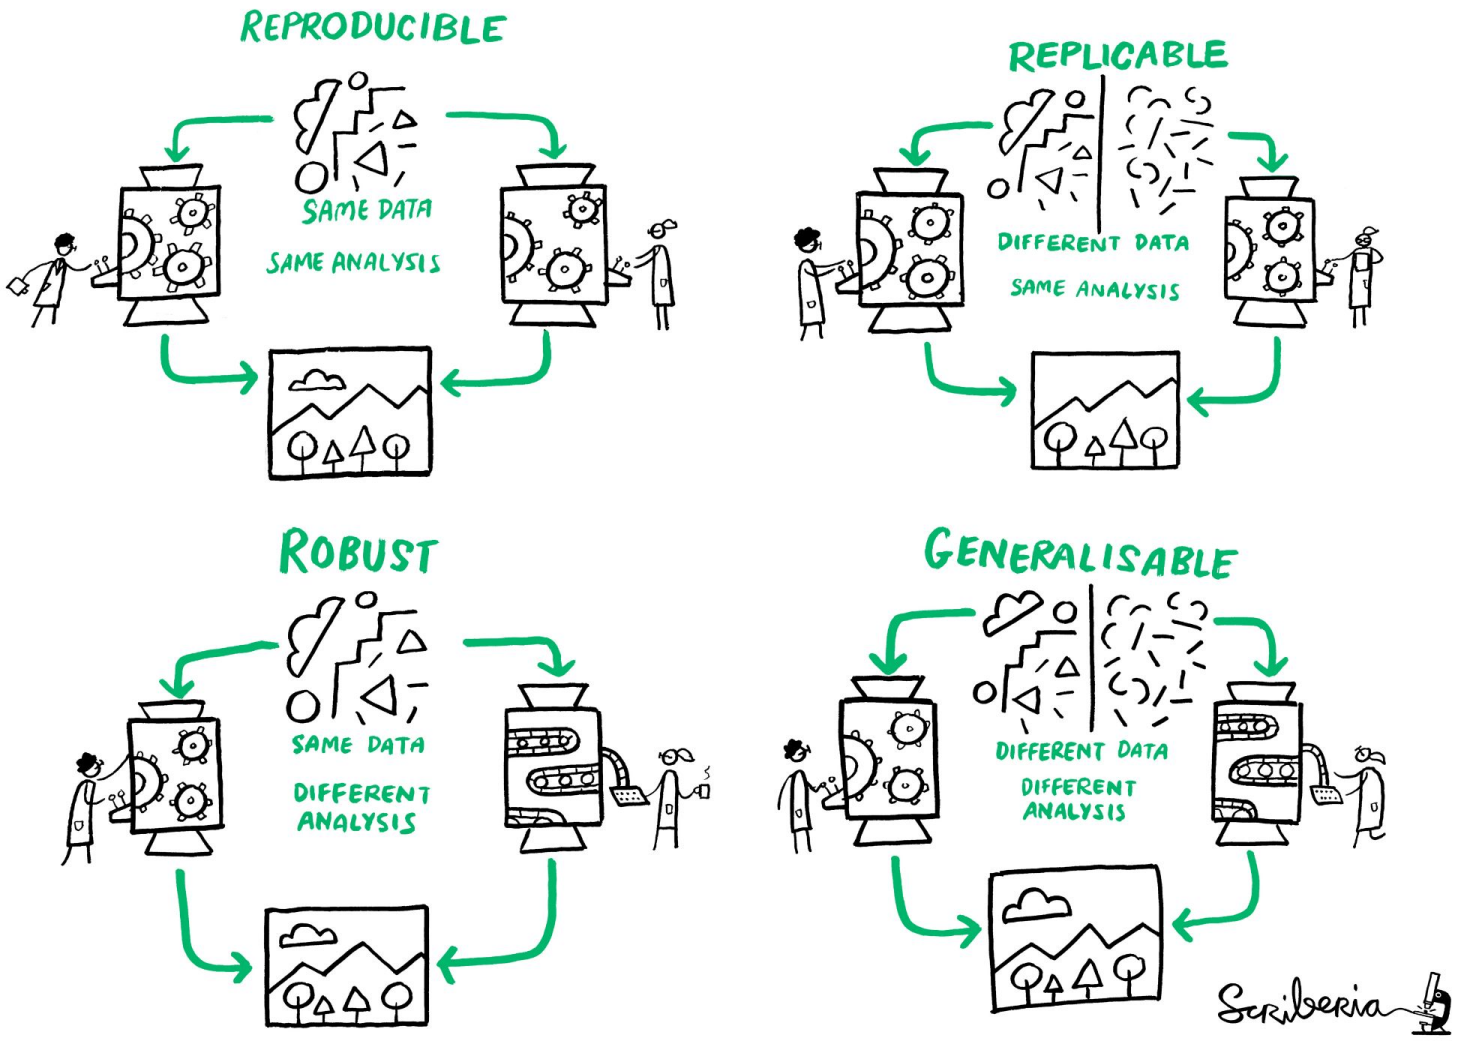

oliviergimenez.github.io/reproducible-science-workshop

Makes your own life easier

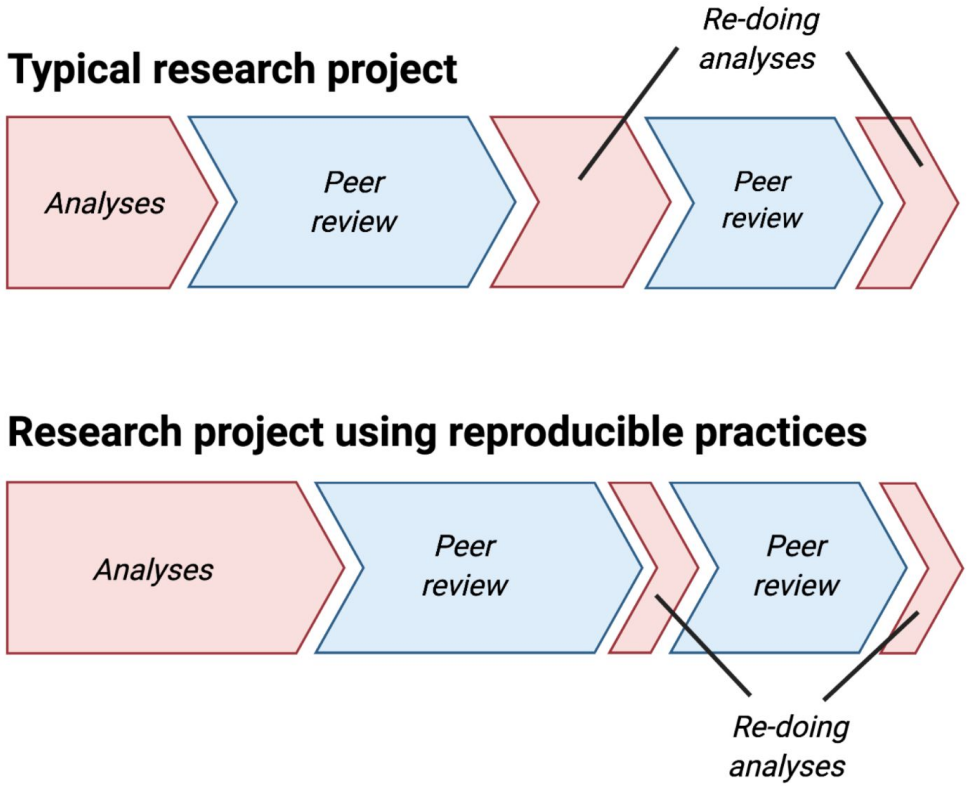

@dsquintana

oliviergimenez.github.io/reproducible-science-workshop

What do we need to do to have reproducible research?

Reproducibility checklist

- Don’t do anything by hand (even “one-off” tasks)

Reproducibility checklist

- Don’t do anything by hand (even “one-off” tasks)
- Script every interaction with data:
  - Data collection
  - Moving data on your computer
  - Formatting datasets
  - Cleaning data
  - Exploratory data analysis
  - Main analyses
  - Report generation

Reproducibility checklist

- Don’t do anything by hand (even “one-off” tasks)
- Script every interaction with data:
  - Data collection
  - Moving data on your computer
  - Formatting datasets
  - Cleaning data
  - Exploratory data analysis
  - Main analyses
  - Report generation
- Minimise interactivity/point and click interactions

Reproducibility checklist

- Don’t do anything by hand (even “one-off” tasks)
- Script every interaction with data:
  - Data collection
  - Moving data on your computer
  - Formatting datasets
  - Cleaning data
  - Exploratory data analysis
  - Main analyses
  - Report generation
- Minimise interactivity/point and click interactions
- Version control all data, code, and documentation

Reproducibility checklist

- Don’t do anything by hand (even “one-off” tasks)
- Script every interaction with data:
  - Data collection
  - Moving data on your computer
  - Formatting datasets
  - Cleaning data
  - Exploratory data analysis
  - Main analyses
  - Report generation
- Minimise interactivity/point and click interactions
- Version control all data, code, and documentation
- Use a random seed

Reproducibility checklist

- Don’t do anything by hand (even “one-off” tasks)
- Script every interaction with data:
  - Data collection
  - Moving data on your computer
  - Formatting datasets
  - Cleaning data
  - Exploratory data analysis
  - Main analyses
  - Report generation
- Minimise interactivity/point and click interactions
- Version control all data, code, and documentation
- Use a random seed
- Keep track of the exact version of every library/program you use

How do we actually do these things?

Choose a language that makes it easy to do most/all of your analysis

Choose a language that makes it easy to do most/all of your analysis

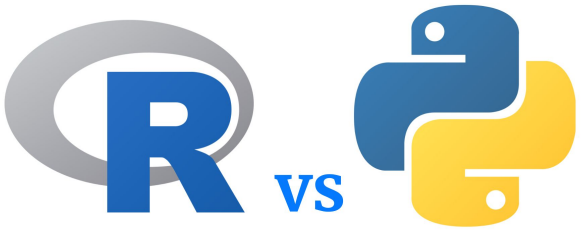

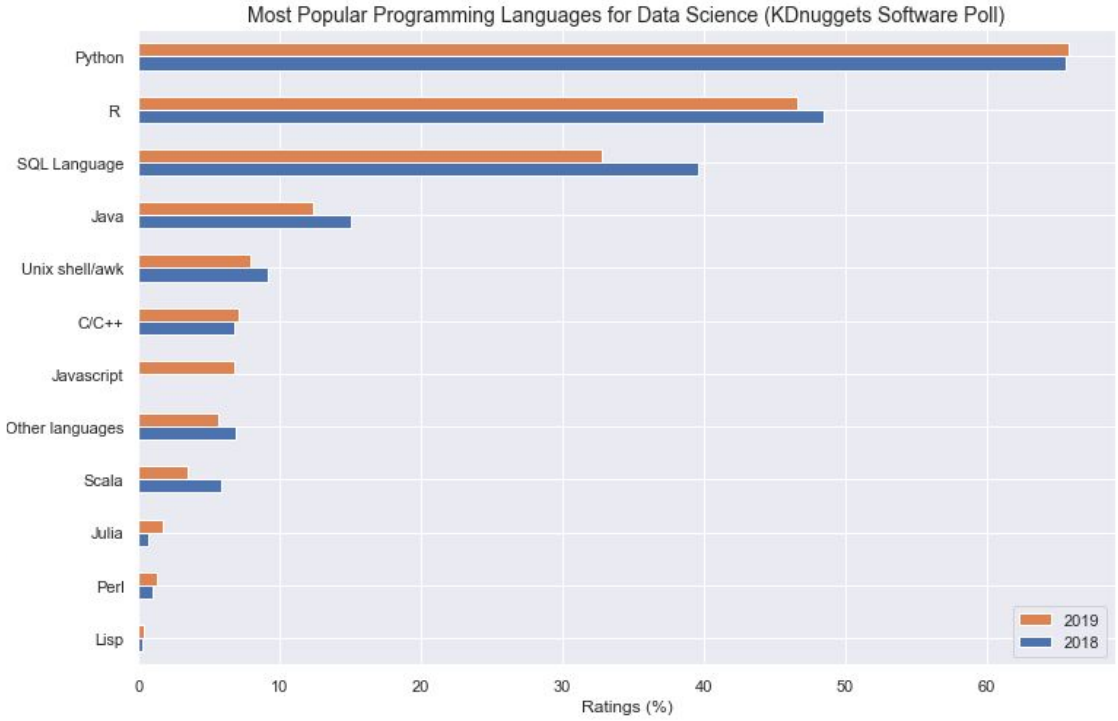

https://www.kdnuggets.com/2019/05/poll-top-data-science-machine-learning-platforms.html

Choose a language that makes it easy to do most/all of your analysis

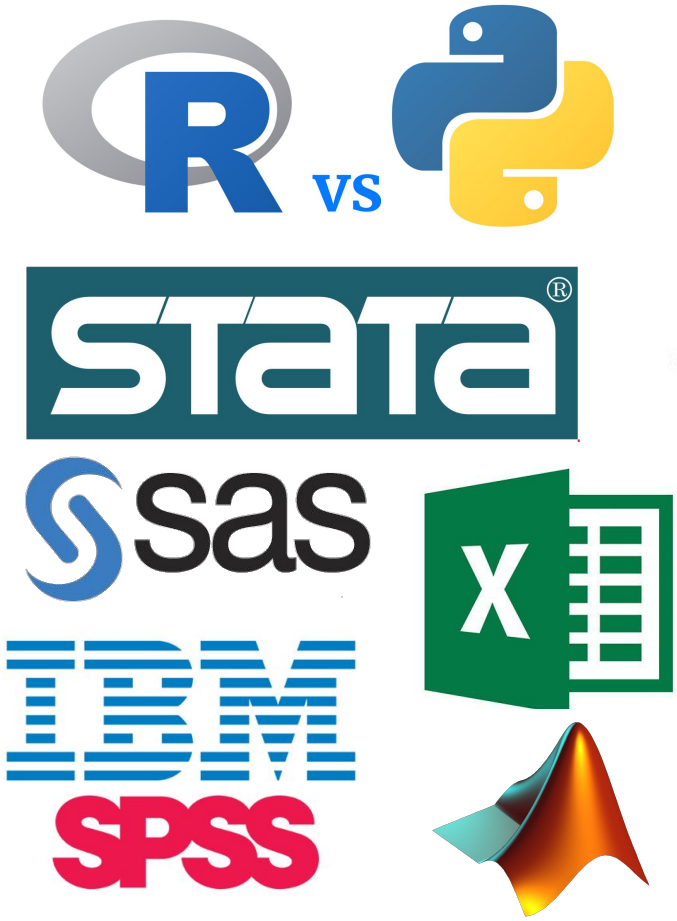

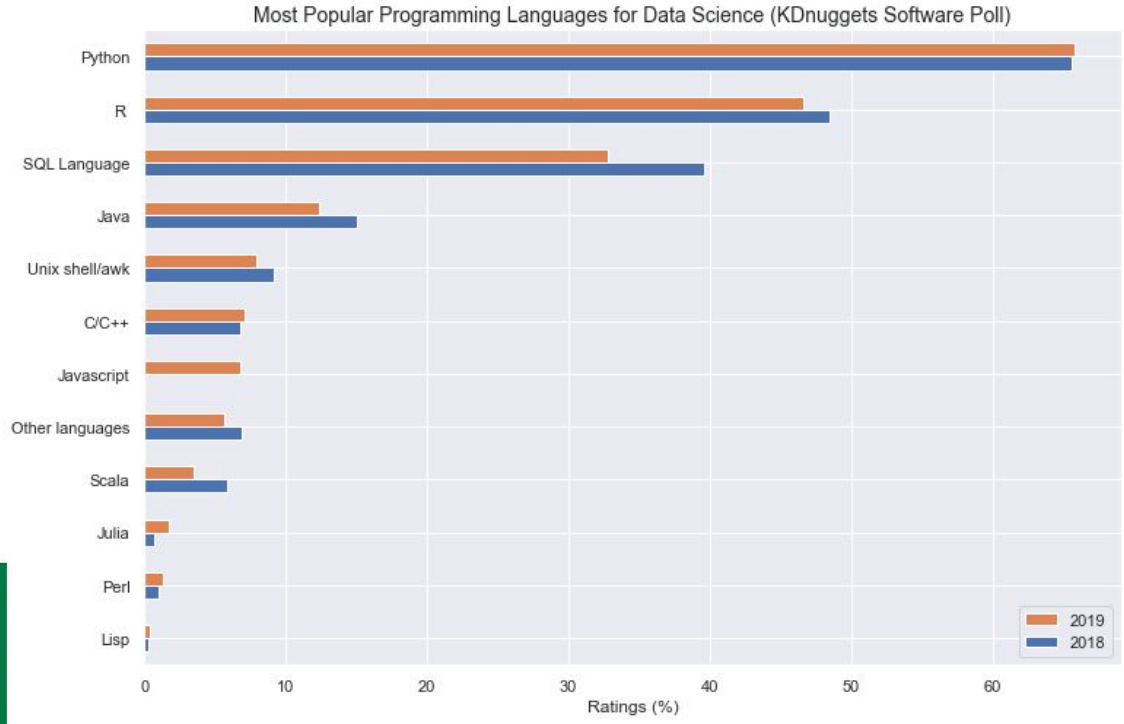

https://www.kdnuggets.com/2019/05/poll-top-data-science-machine-learning-platforms.html

Use a data science focused IDE: Rstudio

set.seed()
sessionInfo()

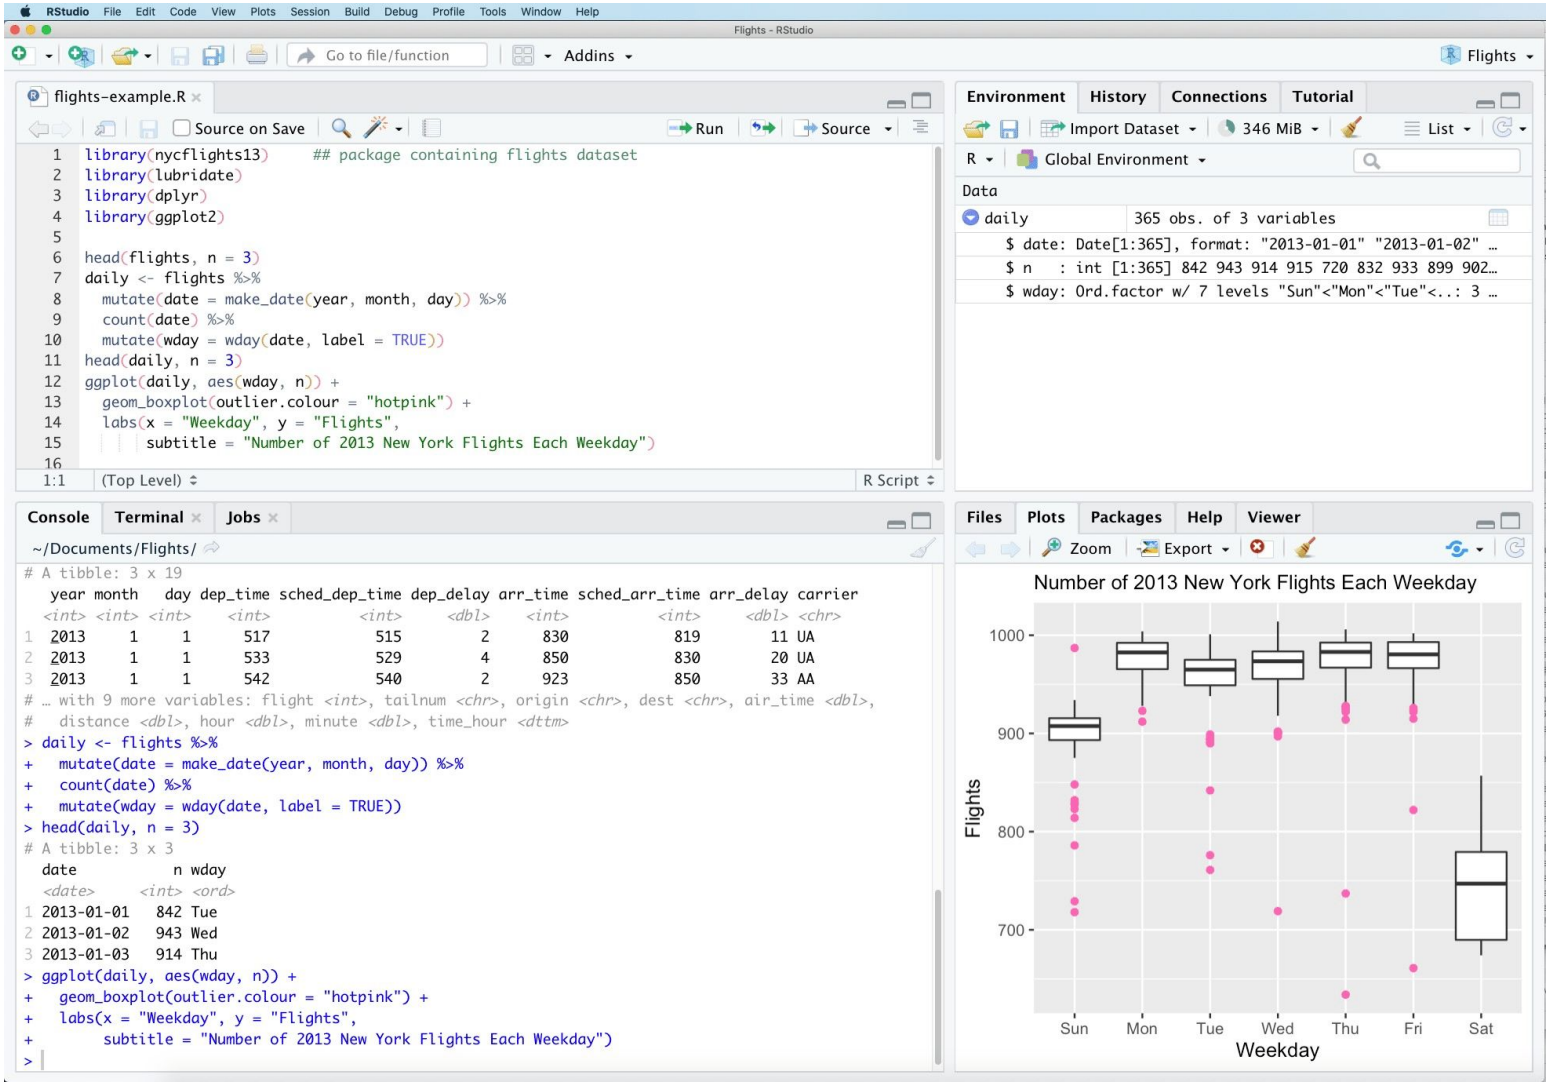

Use notebooks to document analyses: Rmarkdown/Quarto

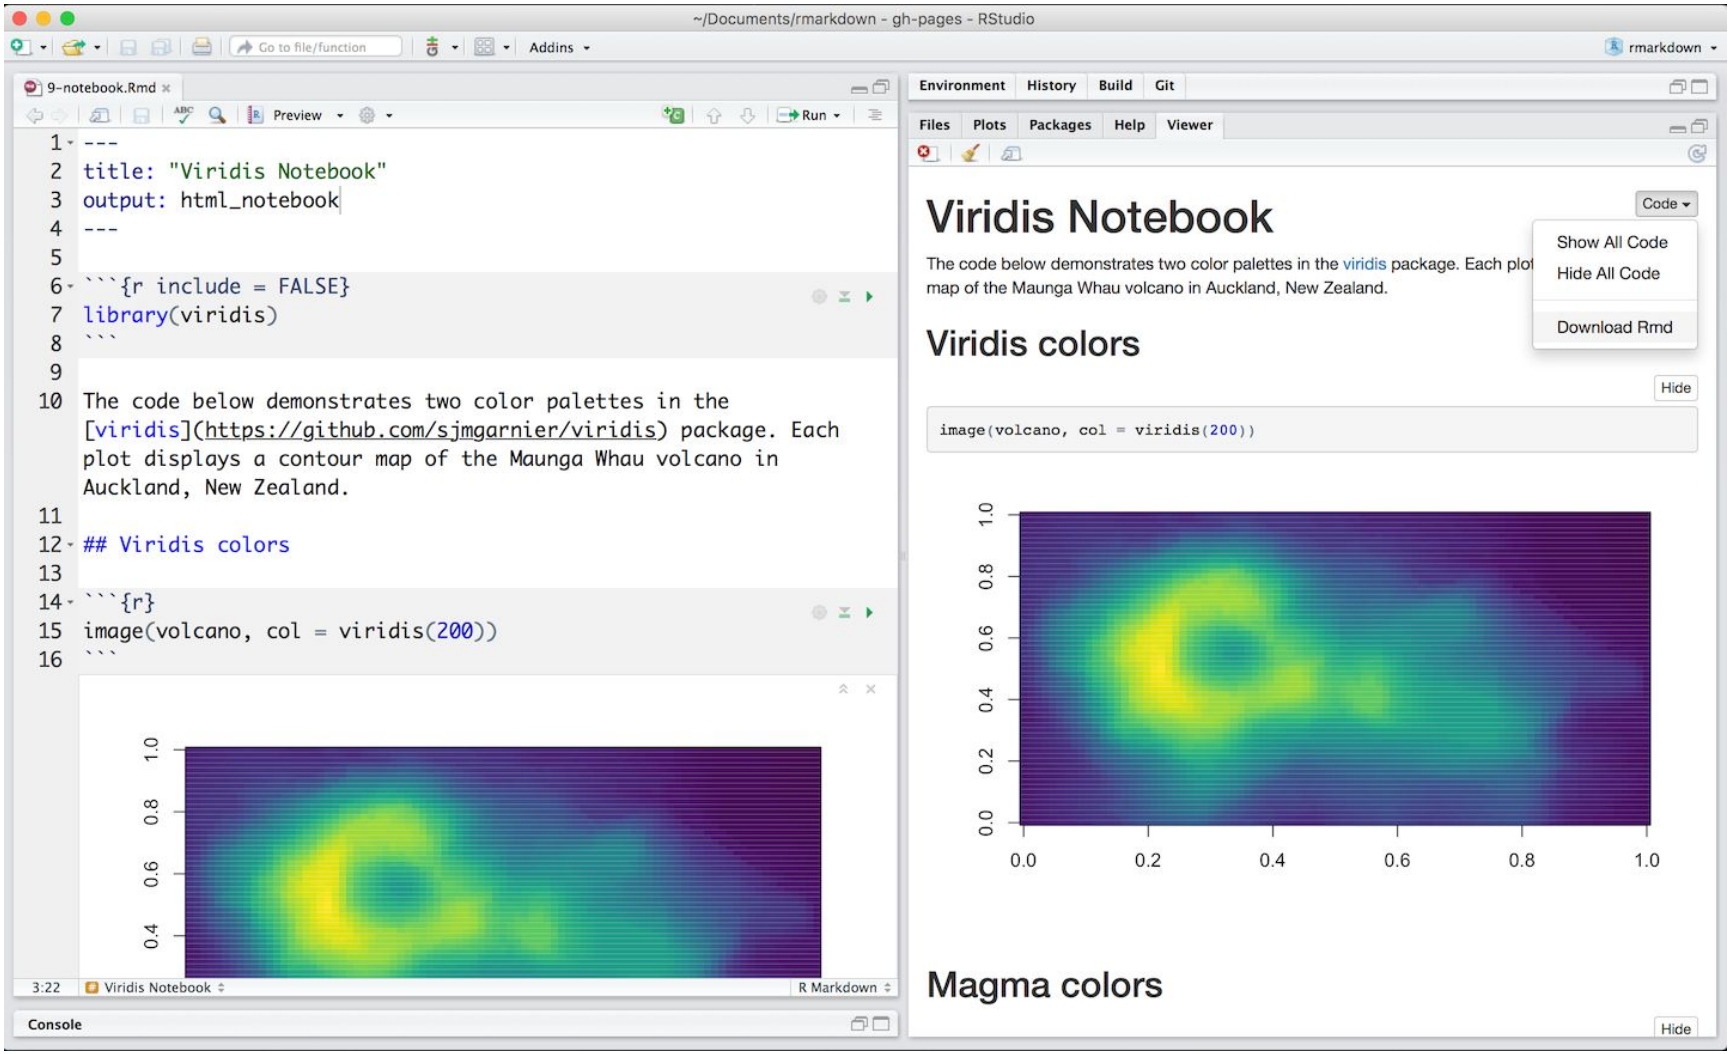

Use notebooks to document analyses: Rmarkdown/Quarto

settings). Therefore, from this time onward, case counts are likely underestimated and the sequenced virus diversity is not necessarily representative of the virus circulating in the overall population.

BC AB SK MB ON QC NS NB NL

Nova Scotia

Additional up-to-date COVID data for this province can be found here:
https://experience.arcgis.com/experience/204d6ed723244dfbb763ca3f913c5cad

Hide

```
plot.variants(region='Nova Scotia')
plot.variants(region='Nova Scotia', scaled=T)
```

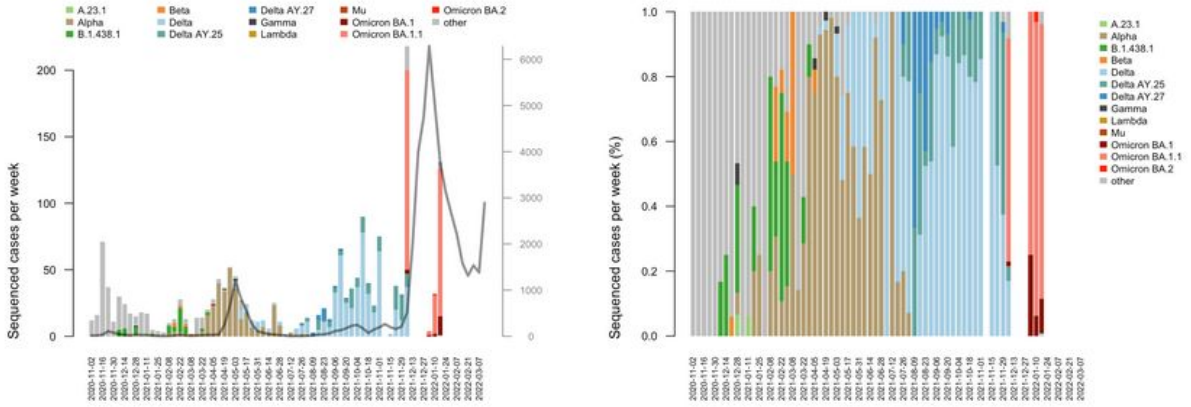

https://covarr-net.github.io/duotang/duotang.html#

Use standard version control systems

- Ever had a nightmare of versioning even when just you?
- Add more people and the chaos grows exponentially!

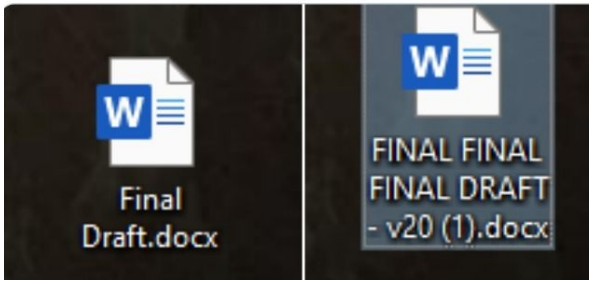

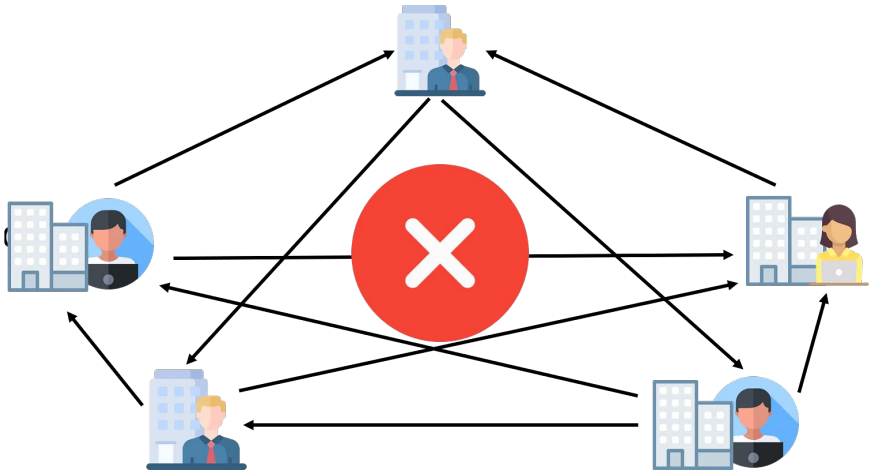

Use standard version control systems

- Ever had a nightmare of versioning even when just you?
- Add more people and the chaos grows exponentially!

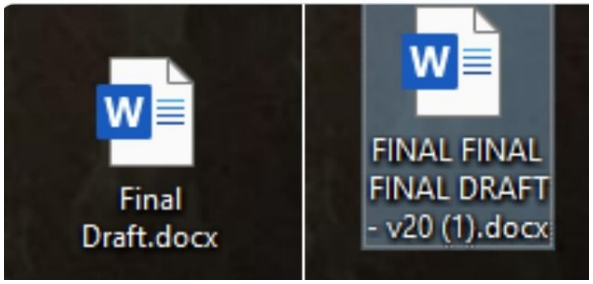

Version control let's you:

- Revert mistakes
- Acts as a comprehensive backup
- Let's you maintain multiple versions of your analysis
- Let's you compare different versions of your co
- Track down the who/what broke the analysis
- Work out why you did something in the past
- Build on someone else's work
- Share your own work
- Experiment without risk

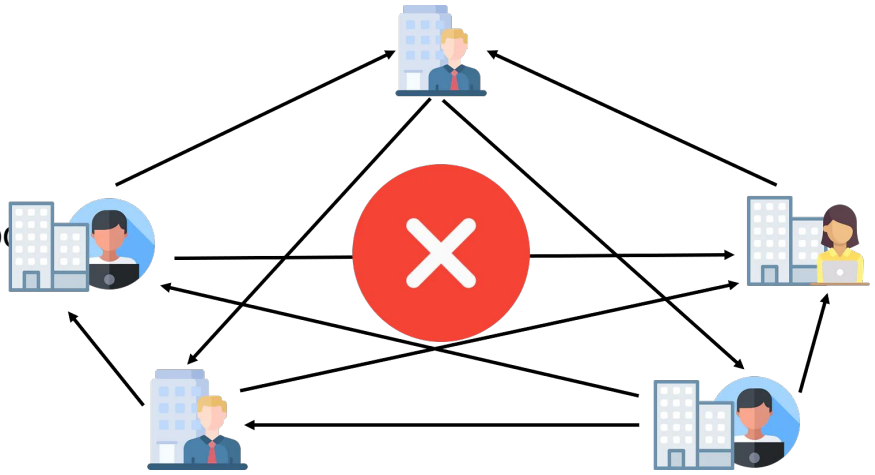

Git Version Control

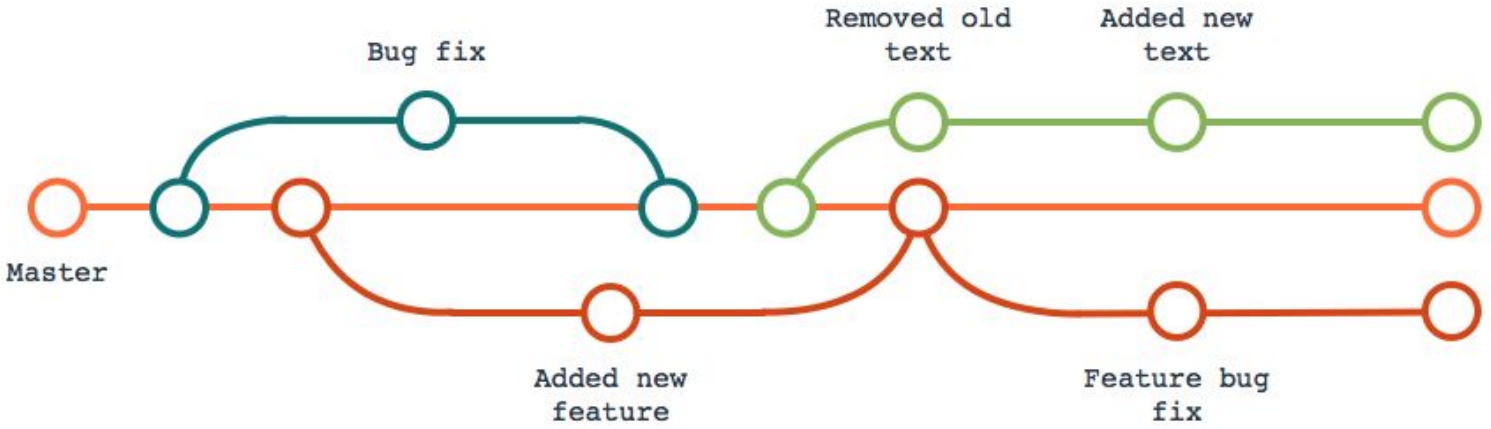

- Most popular
- Decentralised
- Designed for
- GitLab/GitHub Services

Git Version Control

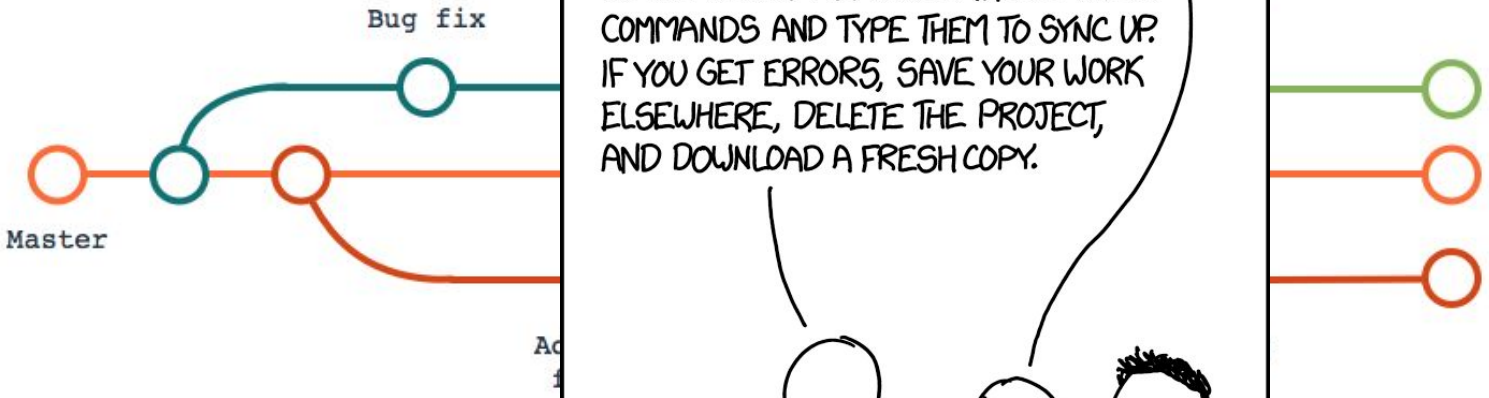

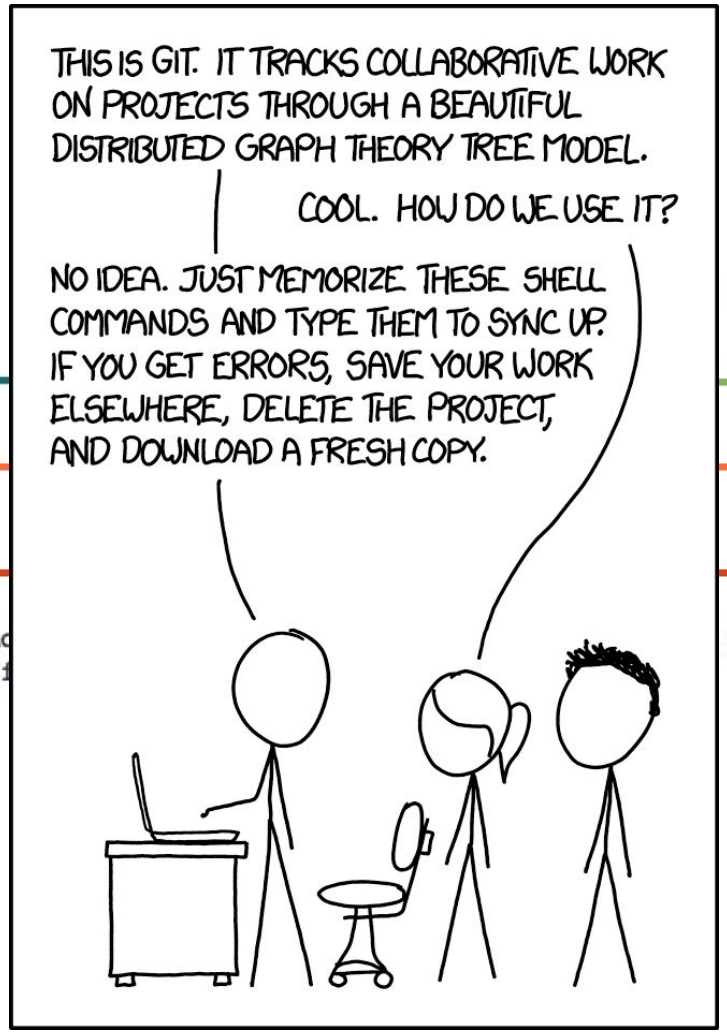

- Most popular
- Decentralised
- Designed for
- GitLab/GitHub Services

Git Workflow

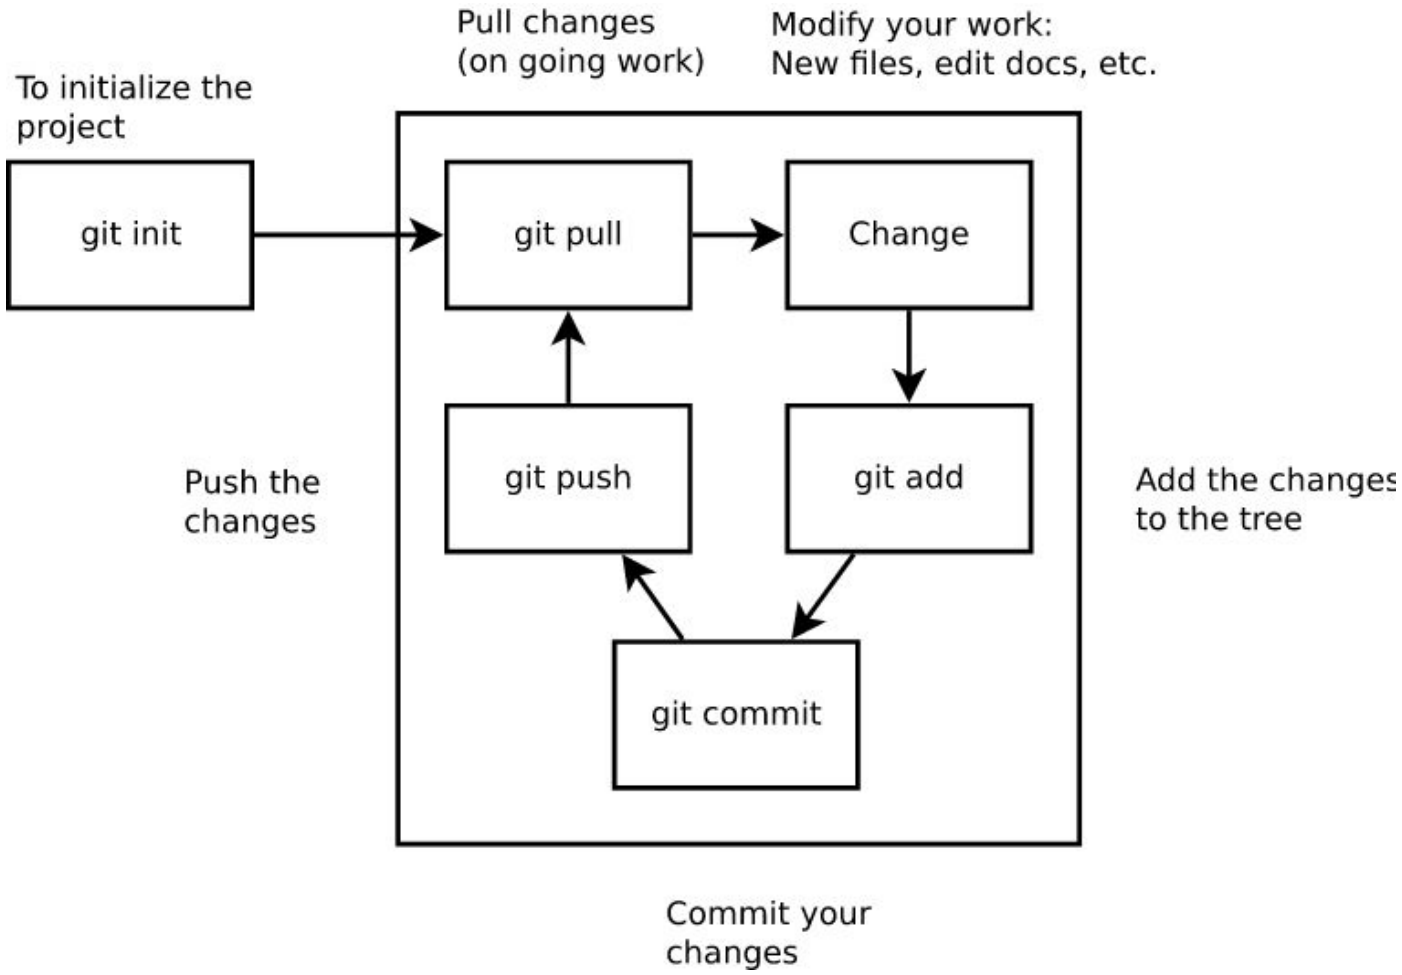

https://uscbiostats.github.io/PM566/slides/02-version-control/slides.html#8

Git is integrated into Rstudio!

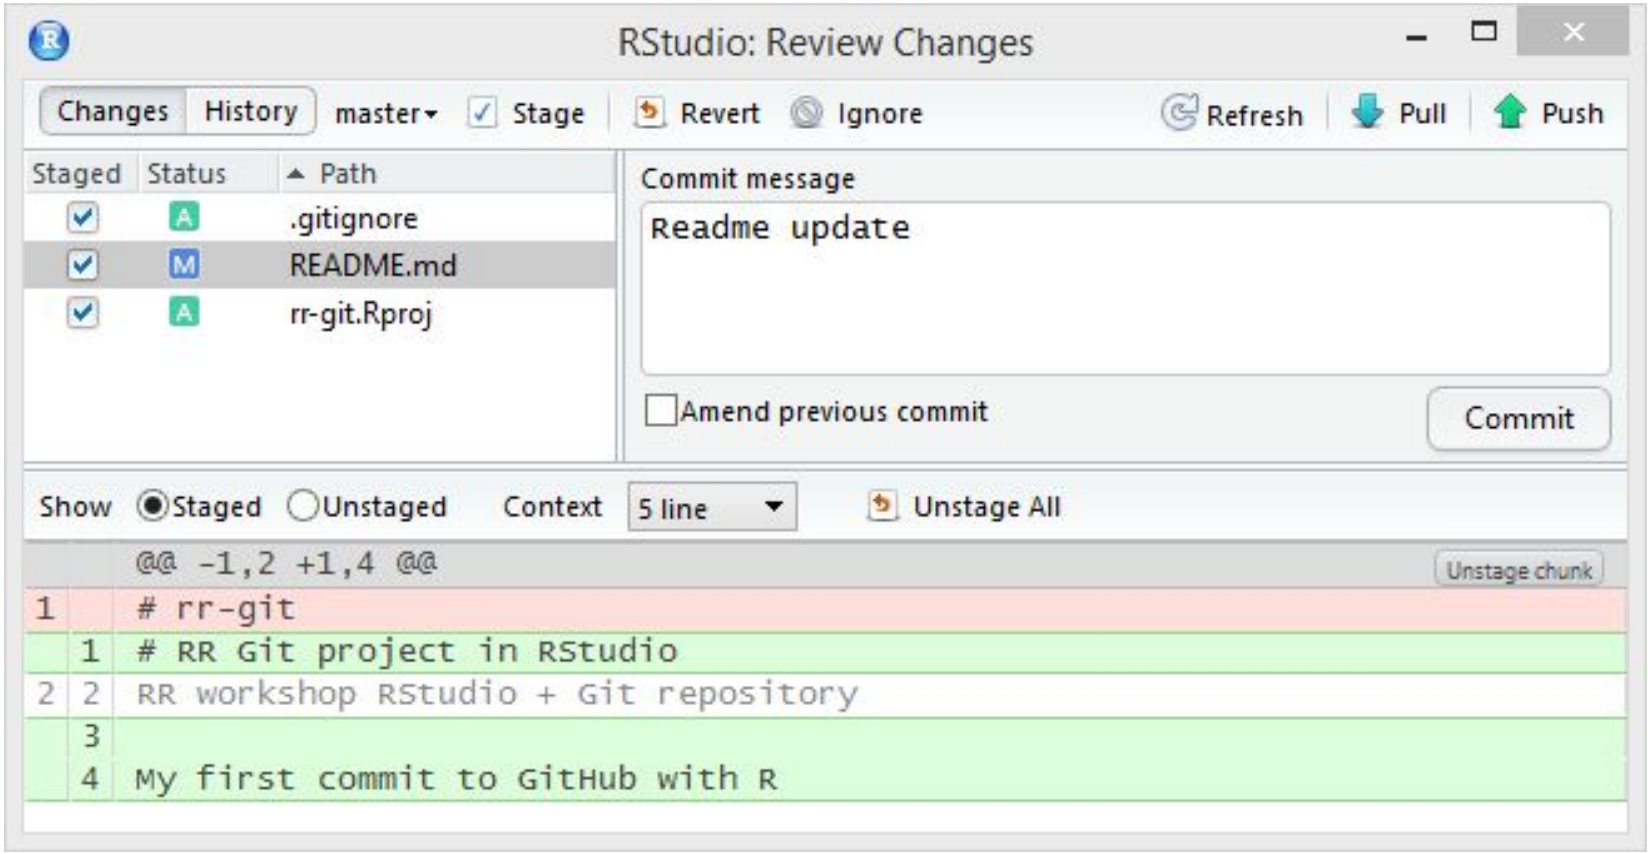

Combine Git+Rmd Notebooks for Reproducibility

1. Add analysis to notebook
2. Add changes to git
3. Find out you made a mistake
4. Revert changes

1. Share notebook with collaborator
2. They make changes
3. You make changes
4. Merge changes into single analysis

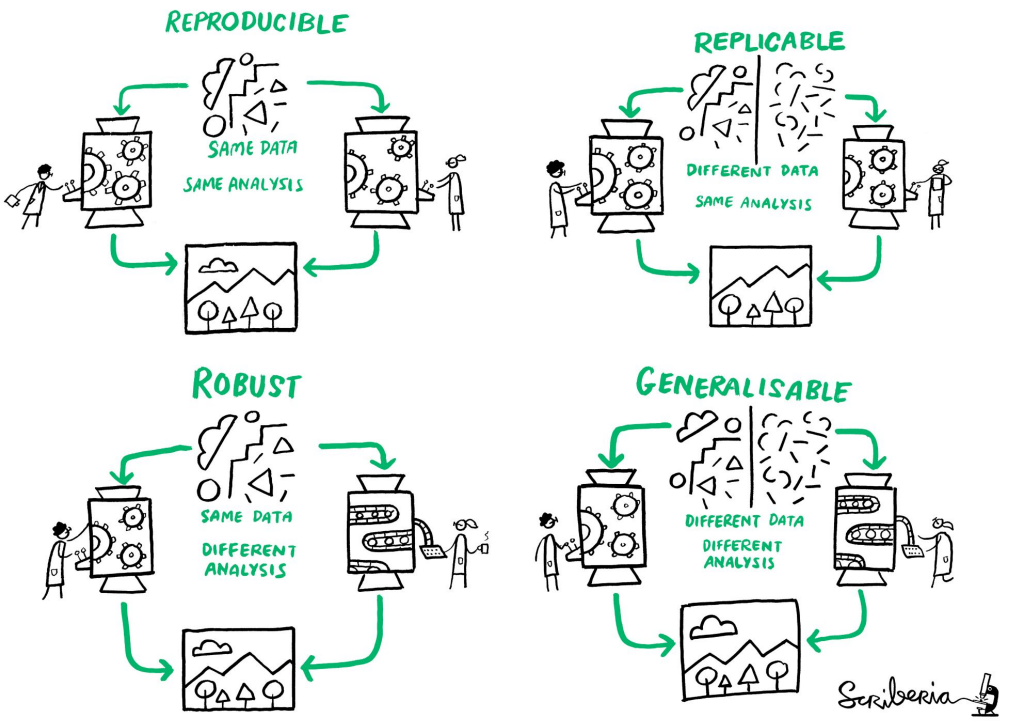

Summary

- Overview of course: Database/EMR/Imaging/Signal
- Main assessments: practicals, journal article presentations, research proposal
- Data science is statistics with an EDA/Inductive/Data-focused Spin
- Health Data Science is a massive and growing area with lots of opportunity and challenges
- R is a powerful and useful tool for health data science
- Reproducibility is vital to good ~~health data~~ science
- Rstudio, Rmarkdown notebooks and Git based version control facilitate that reproducibility

Friday's Practical

- Will go over the practical use of R, Rstudio, Rmd Notebooks, Git
- Try and install rstudio, git, and rmarkdown beforehand.
- 1st practical will not contribute to your course grade

Wednesday's Journal Articles

- **Reproducibility in machine learning for health research: Still a ways to go**

  Matthew B. A. McDermott Shirly Wang Nikki Marinsek Rajesh Ranganath Luca Foschini Marzyeh Ghassemi
  Science Translational Medicine • 24 Mar 2021 • Vol 13, Issue 586 • DOI: 10.1126/scitranslmed.abb1655

- **A Beginner's Guide to Conducting Reproducible Research**

  Jesse M. Alston, Jessica A. Rick First published: 15 January 2021 https://doi.org/10.1002/bes2.1801